



CS471- Parallel Processing

Dr. Ahmed Hesham Mostafa

Lecture 2 – RISC and Instruction Pipelining

Instruction Pipelining

- is a technique for implementing instruction-level parallelism within a single processor. Pipelining attempts to keep every part of the processor busy with some instruction by dividing incoming instructions into a series of sequential steps.
- is a process of arrangement of hardware elements of the CPU such that its overall performance is increased. Simultaneous execution of more than one instruction takes place in a pipelined processor.
- A processor is said to be fully pipelined if it can fetch an instruction on every cycle. Thus, if some instructions or conditions require delays that inhibit fetching new instructions, the processor is not fully pipelined.
- Pipeline is also called Instruction Level Parallelism (ILP).

Pipeline Stages :

- The number of dependent steps varies with the machine architecture. For example:
- The 1956–1961 IBM Stretch project proposed the terms Fetch, Decode, and Execute that have become common.
- The classic RISC pipeline comprises:
 - Instruction fetch
 - Instruction decode and register fetch
 - Execute
 - Memory access
 - Register write back
- The Atmel AVR and the PIC microcontroller each have a two-stage pipeline.

CISC V.S RISC

- **CISC is a complex instruction** , which can perform multiple operations and uses a complex hardware:
- Example MAC R0,R1,R3 which means $R0 = (R1 * R3) + R0$.
- Example ADD [R0],[R1],[R3] which means add the value present in the address in the R1 to the value present in the address in the R3 and store the result in the address present in the R0

CISC V.S RISC programs

- **RISC is a reduced instruction set** , which can only perform simple instructions and uses simple hardware.
- Example ADD R0,R1,R3 which means $R0 = R1 + R3$;
- Example LD R0,[R1] which means move R0 register the value present in the address that present in the R1.

RISC: Reduced Instruction Set Computer

Properties:

- All operations on data apply to data in registers and typically change the entire register (32 or 64 bits per reg);
- Only load and store operations affect memory;
load: move data from mem to reg;
store: move data from reg to mem;
- Only a few instruction formats; all instructions typically being one size.

RISC: Reduced Instruction Set Computer

- RISC has 32 registers and 3 classes of instructions.
- **The first class is ALU instruction.**
- It usually operates on two registers and stores the result in a third register.
- Typical ALU instructions include add, subtract, and logical operations such as AND, OR.

RISC: Reduced Instruction Set Computer

- **The second class is Load and Store instructions** that affect memory.
- They first get the memory address by adding the base register and an offset. The load instruction uses a second register operand as the data destination while the store instruction uses a second register as the data source.
- Load (LD) and store (SD) instructions
operands: base register + offset;
the sum (called *effective address*) is used as a memory address;

RISC: Reduced Instruction Set Computer

- **The third class is branches and jumps** that make conditional transfers of control.
- A branch instruction consists of two phases. The first phase specifies the branch condition to decide whether the branch should take effect.
- It's just like an if-else in programming language.
- The branch condition can be verified via a set of condition bits or comparison between two registers.
- If the branch condition is satisfied, the second phase decides the branch destination. That is, where to fetch the next instruction.

RISC: Reduced Instruction Set Computer

- at most **5** clock cycles per instruction

IF ID EX MEM WB

- **Instruction Fetch cycle**
- The first cycle is IF for fetching the current instruction from memory.
- To do so, the processor sends the program counter to memory, fetches the current instruction from there and then increments the program counter by 4. (as each mem cell 8 bit and instruction is 32 bit so instruction need 4 mem cells)
- $PC = PC + 4;$ //each instr is 4 bytes

RISC: Reduced Instruction Set Computer

- at most **5** clock cycles per instruction

IF ID EX MEM WB

- **Instruction Decode/register fetch cycle**
- decode the instruction;
- read the registers (corresponding to register source specifiers);

RISC: Reduced Instruction Set Computer

- at most **5** clock cycles per instruction

IF ID EX MEM WB

- **Execution/effective address cycle**
- The third cycle is EX for calculating the effective memory address.
- ALU operates on the operands fetched in the ID cycle.
- There are 3 classes of functions according to the instruction type. **The first class is memory reference,** ALU adds base register and offset to form effective address.

RISC: Reduced Instruction Set Computer

- at most **5** clock cycles per instruction

IF ID EX MEM WB

- **Execution/effective address cycle**
- **The second class is register-register ALU instruction:** in this case, ALU performs the operation specified by opcode on the values in the register file.

RISC: Reduced Instruction Set Computer

- at most **5** clock cycles per instruction

IF ID EX MEM WB

- **Execution/effective address cycle**
- **The third class is register-immediate ALU instruction.** For this type, ALU operates on the first value read from the register file and the sign-extended immediate.

RISC: Reduced Instruction Set Computer

- at most **5** clock cycles per instruction

IF ID EX MEM WB

- **MEMory access**
- The fourth cycle is MEM for memory access. there are two types of instructions affecting memory.
- For load instruction, the memory does a read using the effective address.
- For store instruction, the memory writes some data using the effective address.

RISC: Reduced Instruction Set Computer

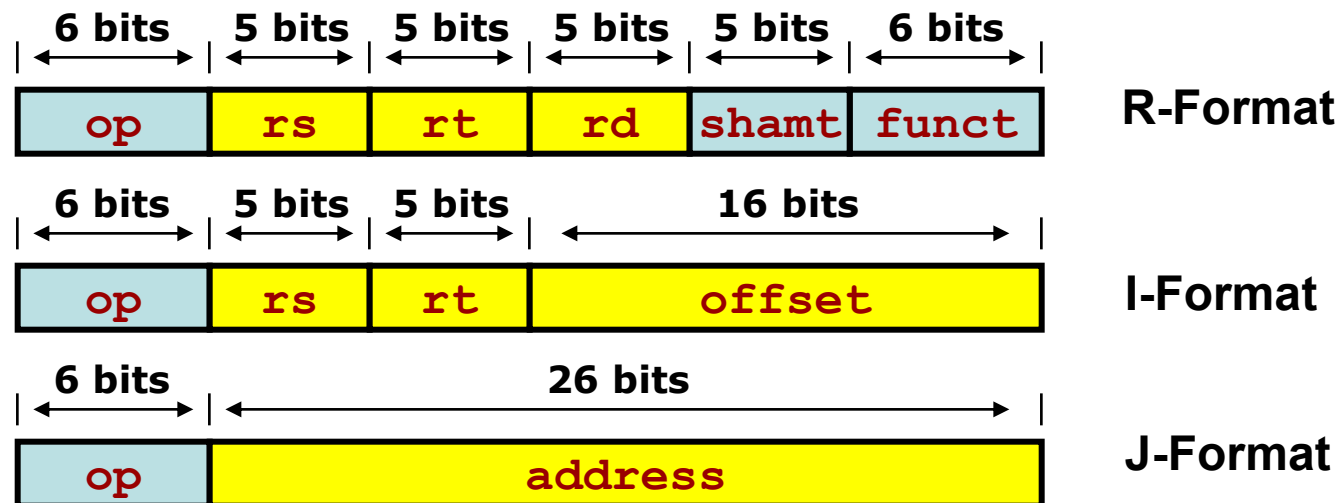
- at most **5** clock cycles per instruction

IF ID EX MEM WB

- **Write-Back cycle**
- The fifth cycle is WB for writing the result into the register file.
- The result may come from the memory if it's a load instruction or from the ALU if it's an ALU instruction.

RISC Instruction Format

- Instructions are divided into three types: R (register), I (immediate), and J (jump).
- Every instruction starts with a 6-bit opcode.
- In addition to the opcode, R-type instructions specify three registers, a shift amount field, and a function field;
- I-type instructions specify two registers and a 16-bit immediate value; J-type instructions follow the opcode with a 26-bit jump target.



Set of Instructions

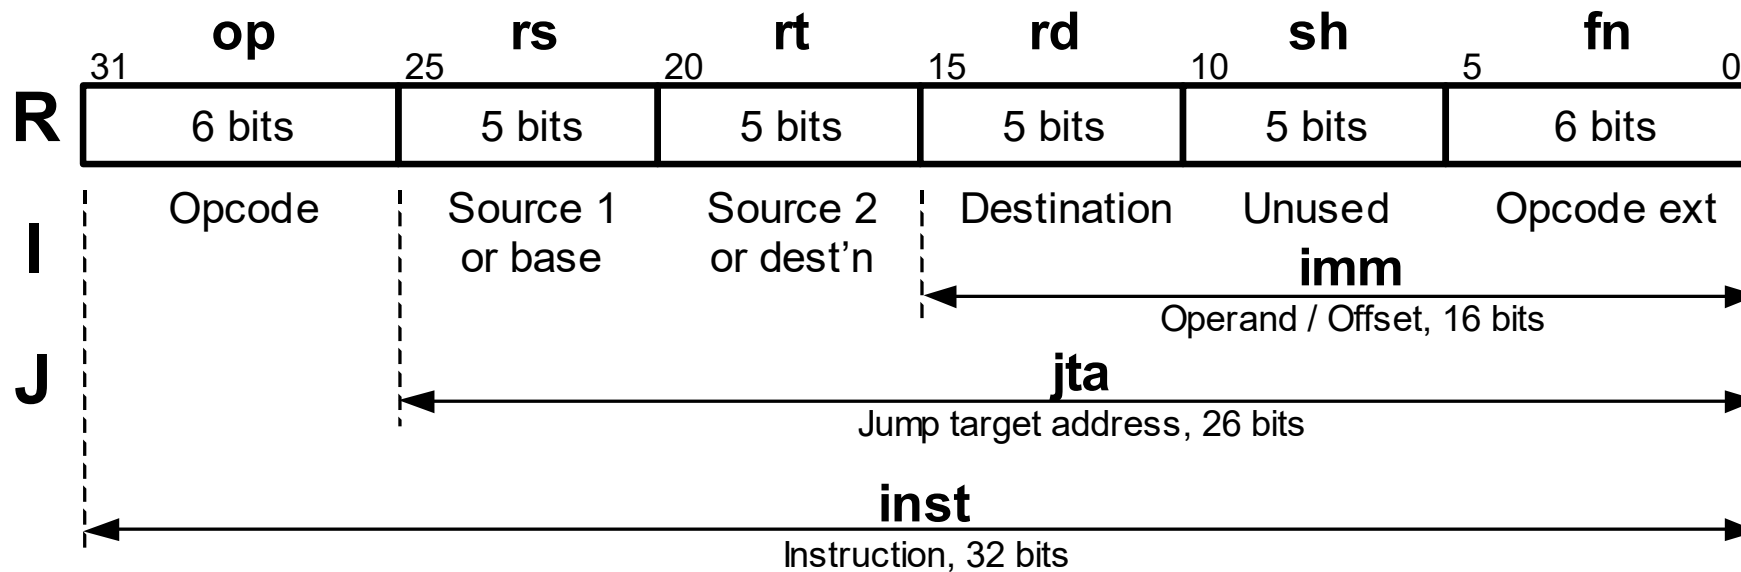


Fig. 13.1 MicroMIPS instruction formats and naming of the various fields.

We will refer to this diagram later

Seven R-format ALU instructions (add, sub, slt, and, or, xor, nor)

Six I-format ALU instructions (lui, addi, slti, andi, ori, xori)

Two I-format memory access instructions (lw, sw)

Three I-format conditional branch instructions (bltz, beq, bne)

Four unconditional jump instructions (j, jr, jal, syscall)

The MicroMIPS Instruction Set

Including
`la rt,imm(rs)`
(load address)
makes it easier
to write useful
programs

Copy

Arithmetic

Logic

Memory access

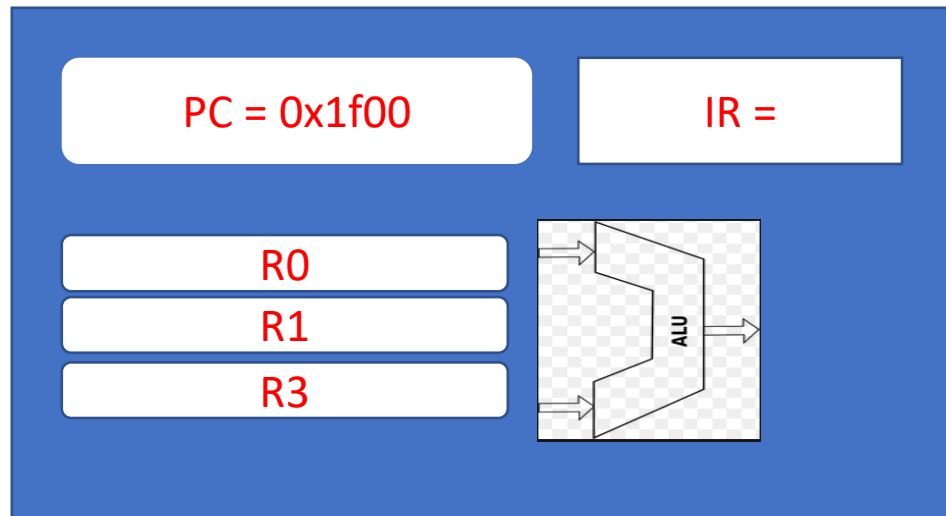
Control transfer

Table 13.1

Instruction	Usage	op	fn
Load upper immediate	<code>lui rt,imm</code>	15	
Add	<code>add rd,rs,rt</code>	0	32
Subtract	<code>sub rd,rs,rt</code>	0	34
Set less than	<code>slt rd,rs,rt</code>	0	42
Add immediate	<code>addi rt,rs,imm</code>	8	
Set less than immediate	<code>slti rd,rs,imm</code>	10	
AND	<code>and rd,rs,rt</code>	0	36
OR	<code>or rd,rs,rt</code>	0	37
XOR	<code>xor rd,rs,rt</code>	0	38
NOR	<code>nor rd,rs,rt</code>	0	39
AND immediate	<code>andi rt,rs,imm</code>	12	
OR immediate	<code>ori rt,rs,imm</code>	13	
XOR immediate	<code>xori rt,rs,imm</code>	14	
Load word	<code>lw rt,imm(rs)</code>	35	
Store word	<code>sw rt,imm(rs)</code>	43	
Jump	<code>j L</code>	2	
Jump register	<code>jr rs</code>	0	8
Branch less than 0	<code>bltz rs,L</code>	1	
Branch equal	<code>beq rs,rt,L</code>	4	
Branch not equal	<code>bne rs,rt,L</code>	5	
Jump and link	<code>jal L</code>	3	
System call	<code>syscall</code>	0	12

A 5-Stage Pipeline (ADD instruction)

ADD R0, R1,R3



CPU

Program Memory

ADD R0, R1,R3

0x0056 0ab0f

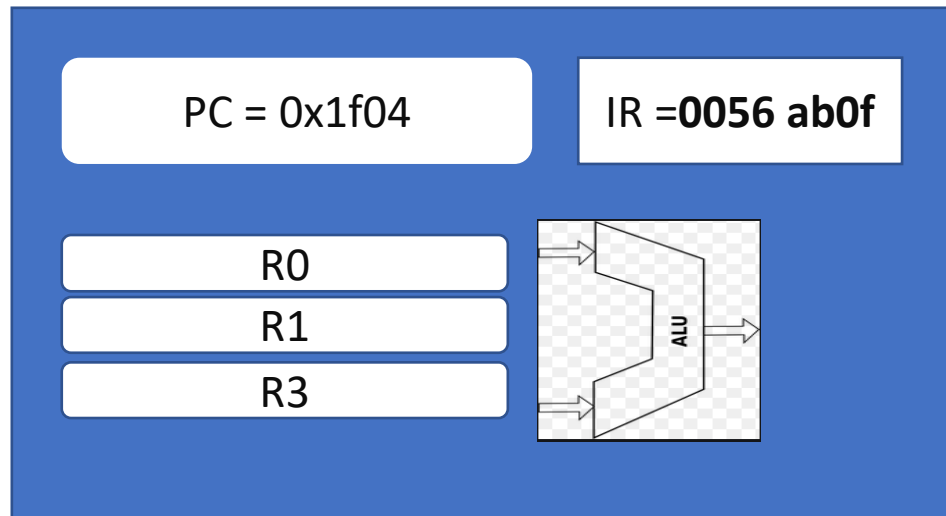
0x01f00

Data Memory

0x 1224 0a00

A 5-Stage Pipeline (ADD instruction)

ADD R0, R1,R3



CPU

Program Memory

ADD R0, R1,R3

0x1856 0000

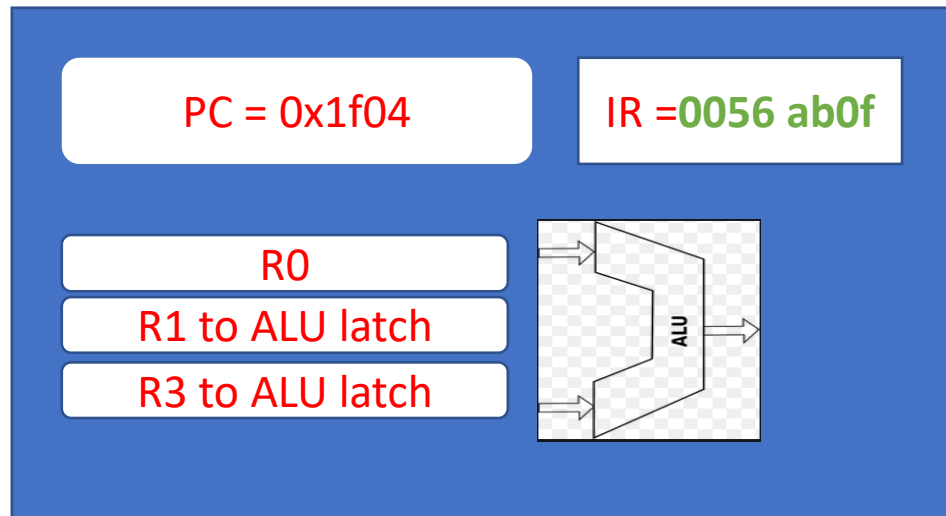
0x01f00

Data Memory

0x 1224 0a00

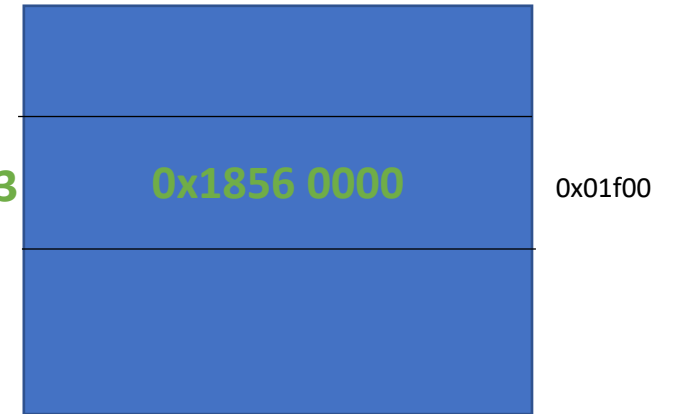
A 5-Stage Pipeline (ADD instruction)

ADD R0, R1,R3



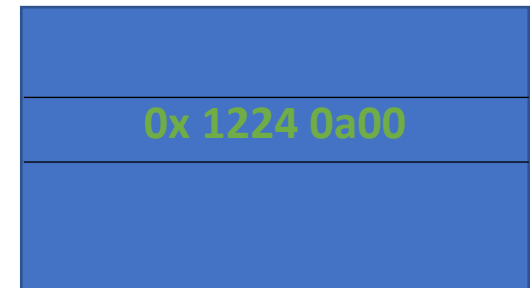
CPU

Program Memory



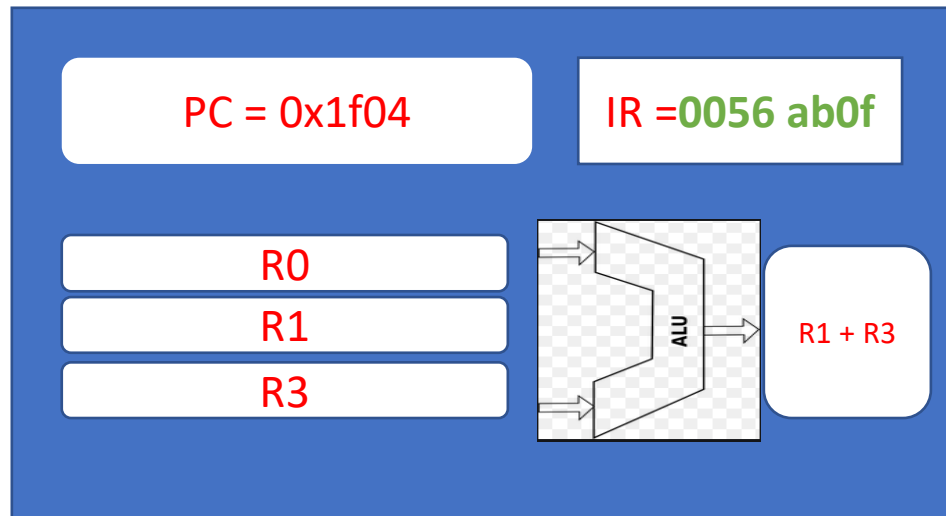
ADD R0, R1,R3

Data Memory



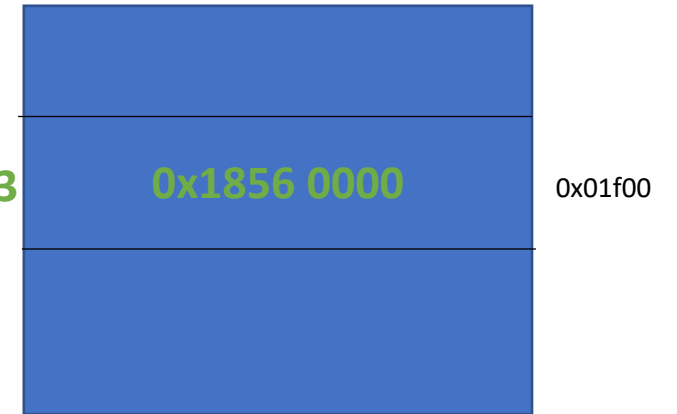
A 5-Stage Pipeline (ADD instruction)

ADD R0, R1,R3



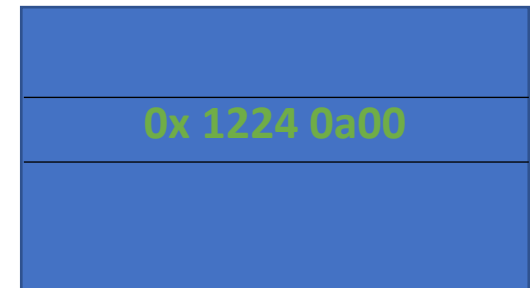
CPU

Program Memory



ADD R0, R1,R3

Data Memory

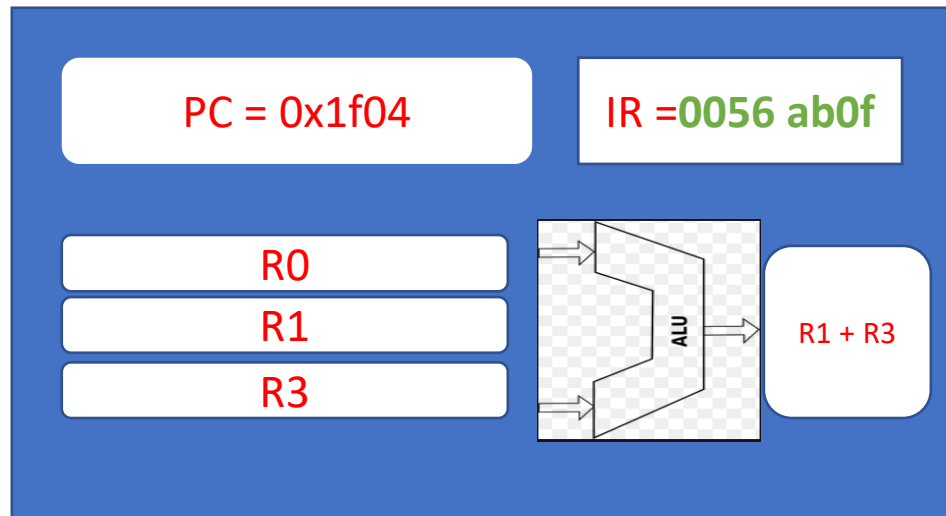


A 5-Stage Pipeline (ADD instruction)

ADD R0, R1,R3

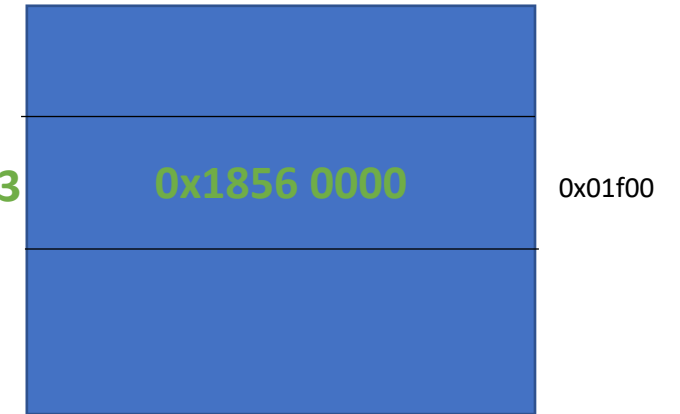


Memory access stage is a dummy state in the Add instruction

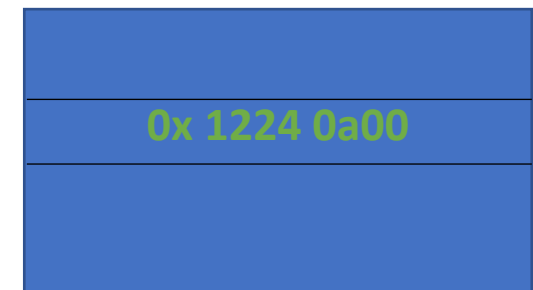


CPU

Program Memory

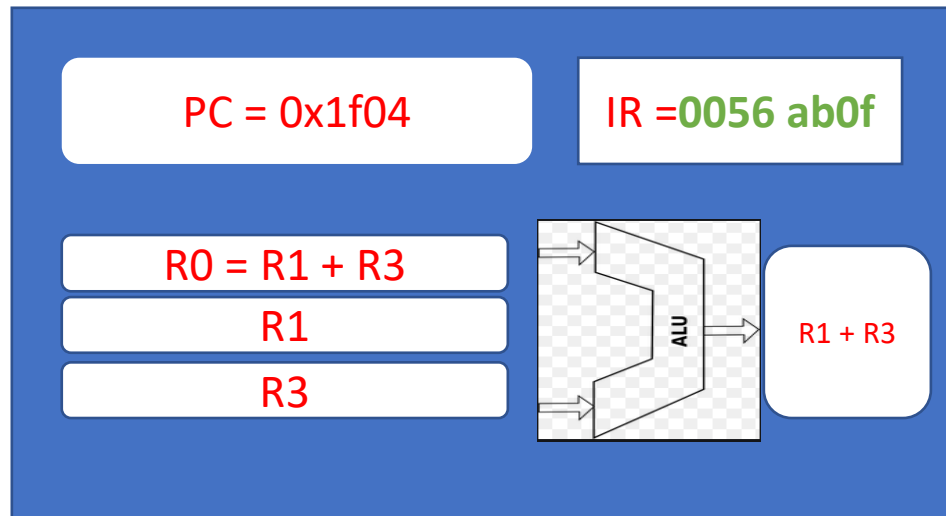


Data Memory



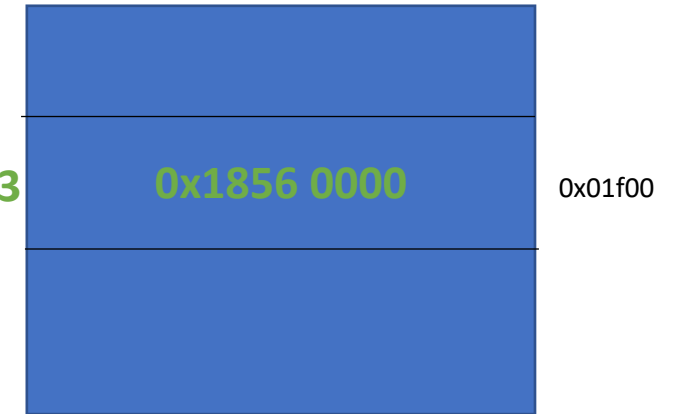
A 5-Stage Pipeline (ADD instruction)

ADD R0, R1, R3

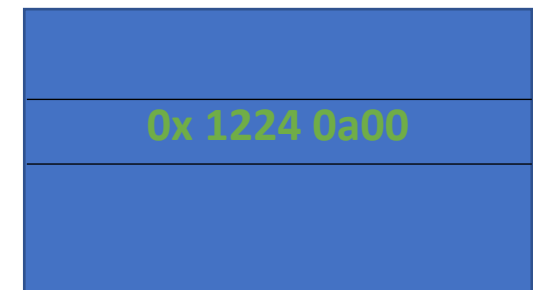


CPU

Program Memory

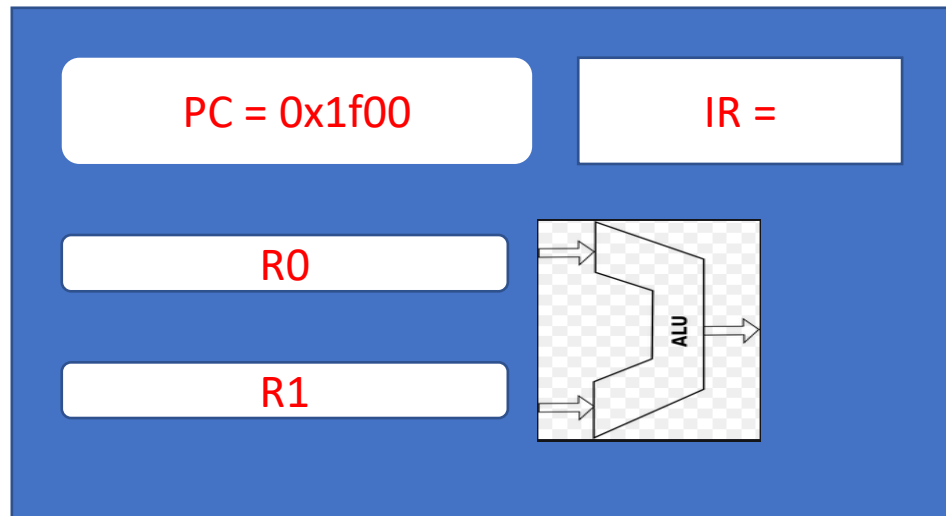


Data Memory



A 5-Stage Pipeline (LD instruction)

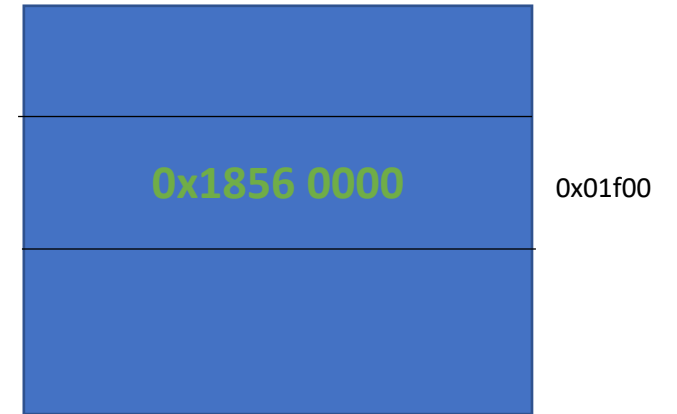
LD R1,24[R0]



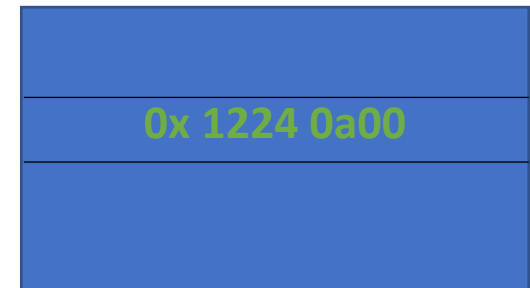
CPU

LD R1,24[R0]

Program Memory

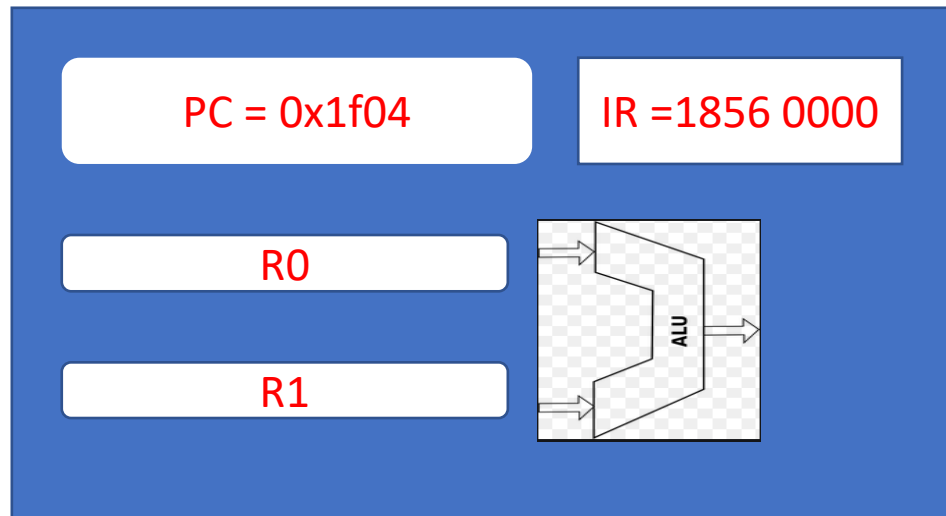


Data Memory



A 5-Stage Pipeline (LD instruction)

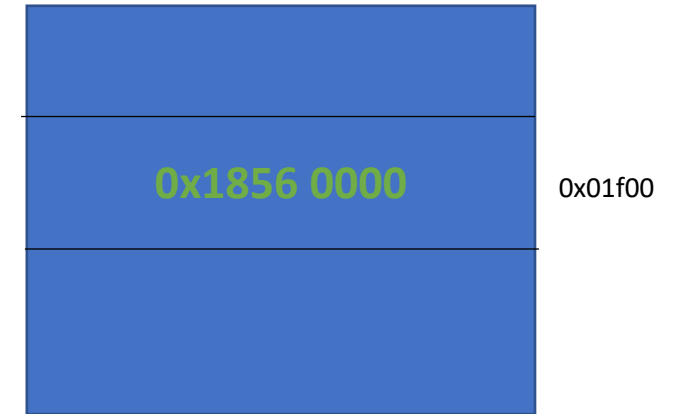
LD R1,24[R0]



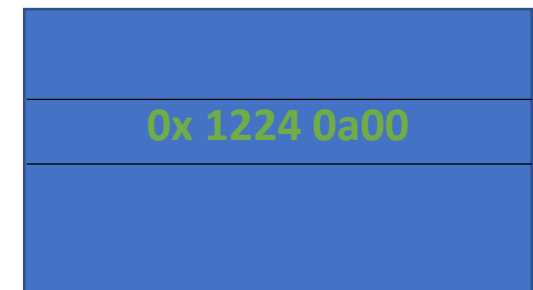
CPU

LD R1,24[R0]

Program Memory

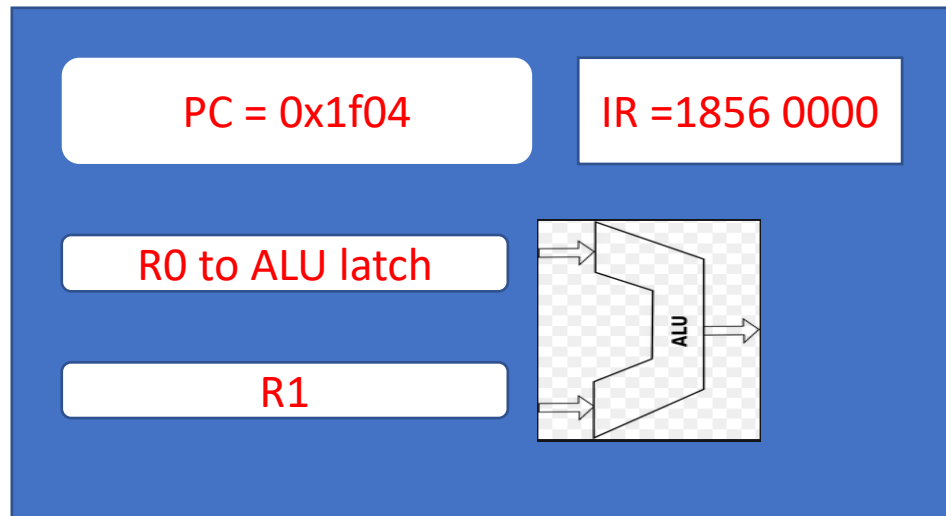


Data Memory



A 5-Stage Pipeline (LD instruction)

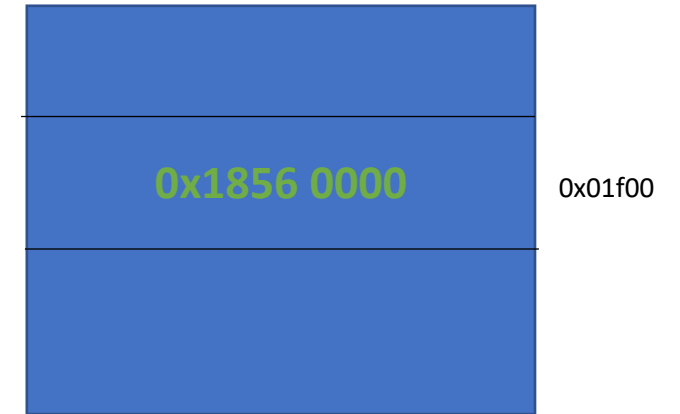
LD R1,24[R0]



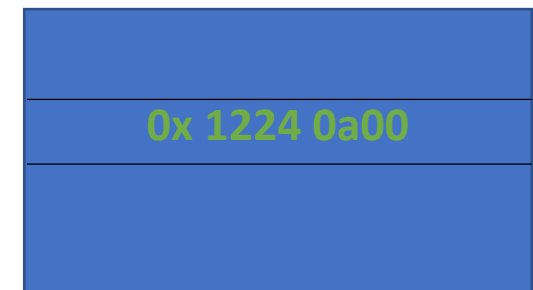
CPU

LD R1,24[R0]

Program Memory

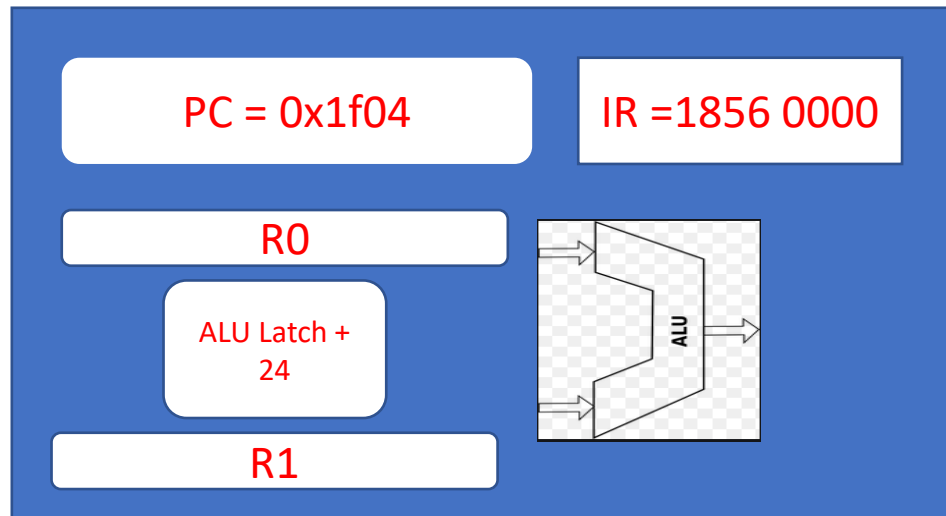


Data Memory



A 5-Stage Pipeline (LD instruction)

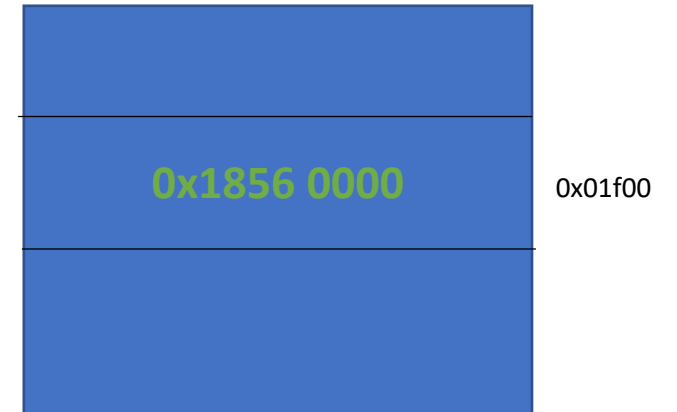
LD R1,24[R0]



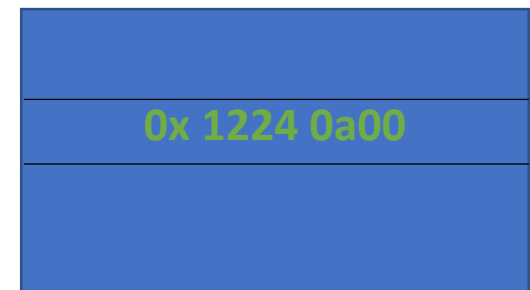
CPU

LD R1,24[R0]

Program Memory

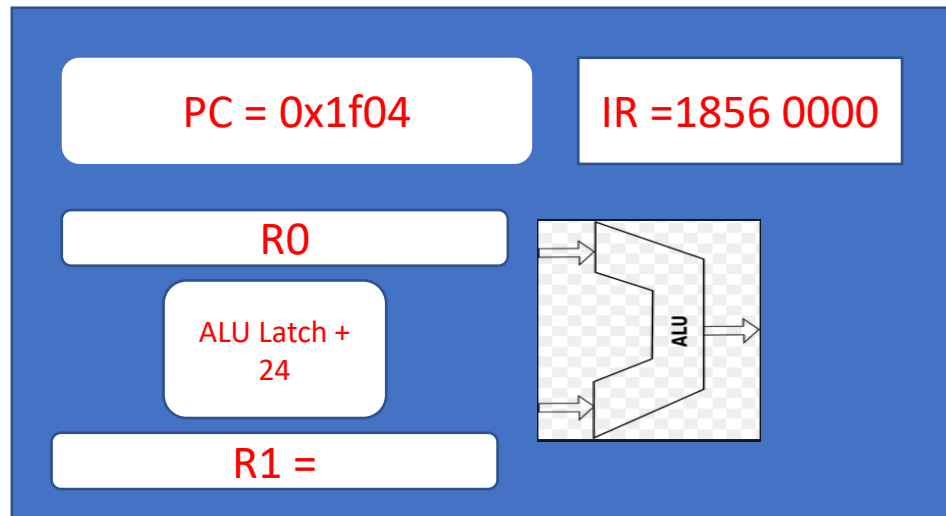


Data Memory



A 5-Stage Pipeline (LD instruction)

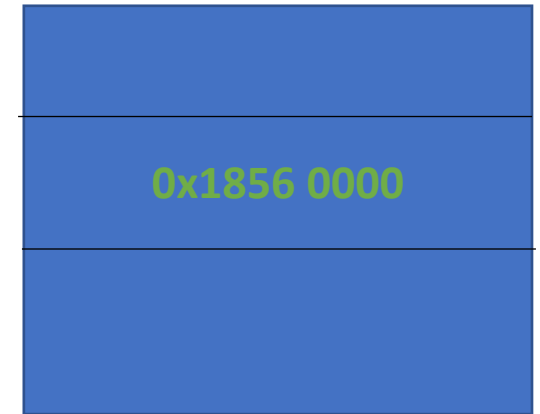
LD R1,24[R0]



CPU

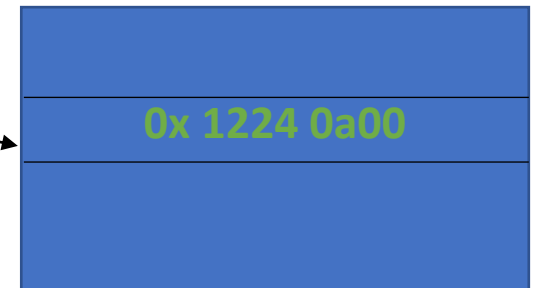
LD R1,24[R0]

Program Memory



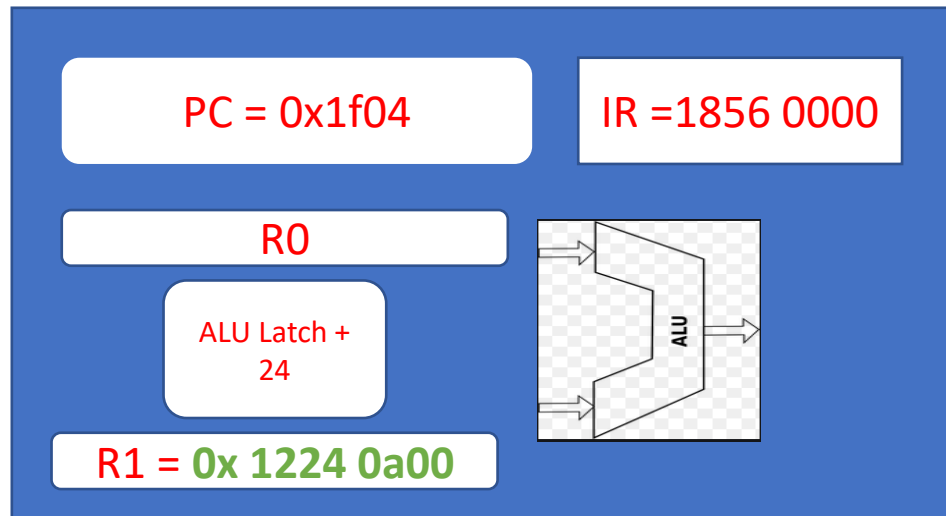
0x01f00

Data Memory



A 5-Stage Pipeline (LD instruction)

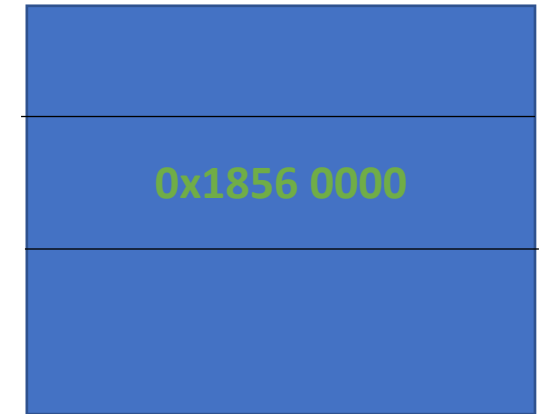
LD R1,24[R0]



CPU

LD R1,24[R0]

Program Memory



0x01f00

Data Memory



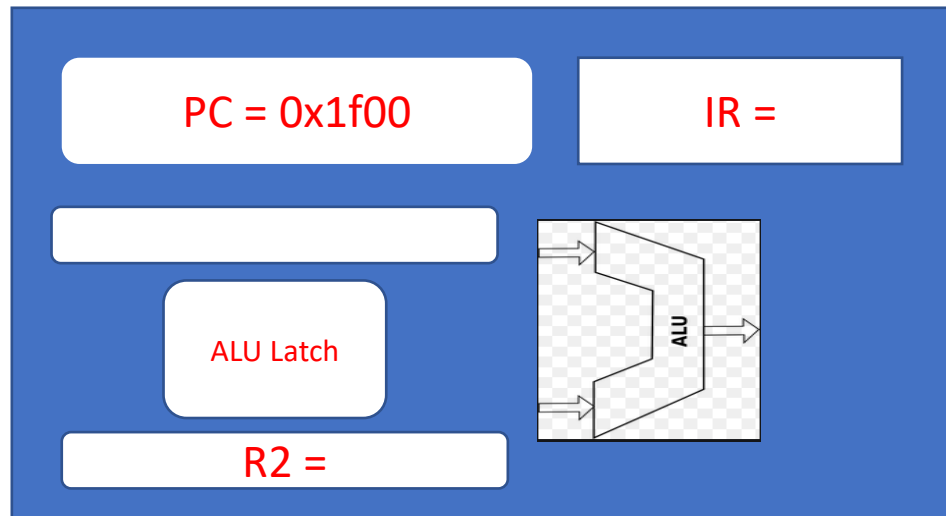
A 5-Stage Pipeline (Branch instruction)

- If the instruction decoded was a branch or jump, the target address of the branch or jump was computed in parallel with reading the register file.
- The branch condition is computed after the register file is read, and if the branch is taken or if the instruction is a jump, the PC predictor in the first stage is assigned the branch target, rather than the incremented PC that has been computed.
- So the branch instruction and call instruction are completed in two cycles.

A 5-Stage Pipeline (Branch instruction)

BEZ R2,100; // If R2 is zero branch to location PC+100.

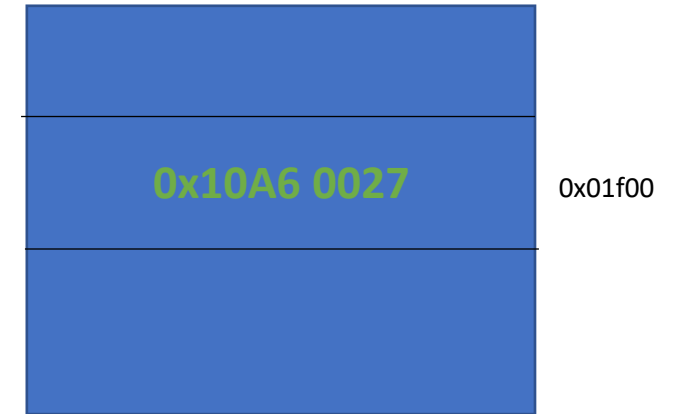
IF	ID	Execute	Memory	Store
----	----	---------	--------	-------



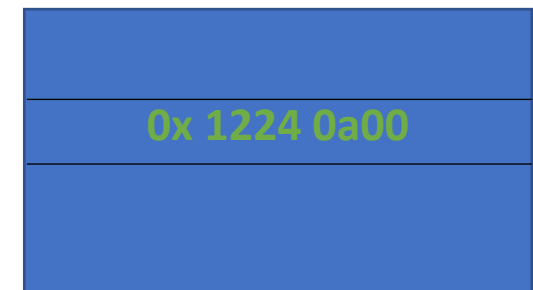
CPU

BEZ R2,100;

Program Memory

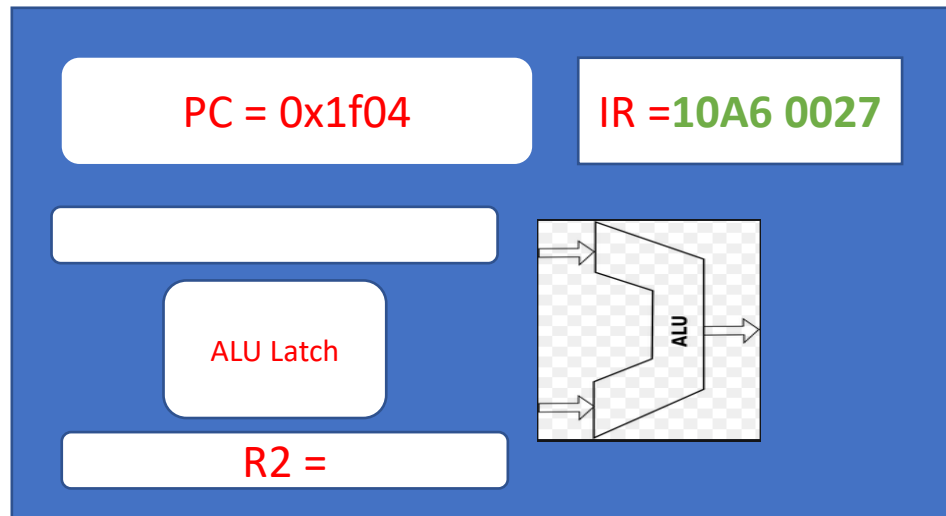


Data Memory



A 5-Stage Pipeline (Branch instruction)

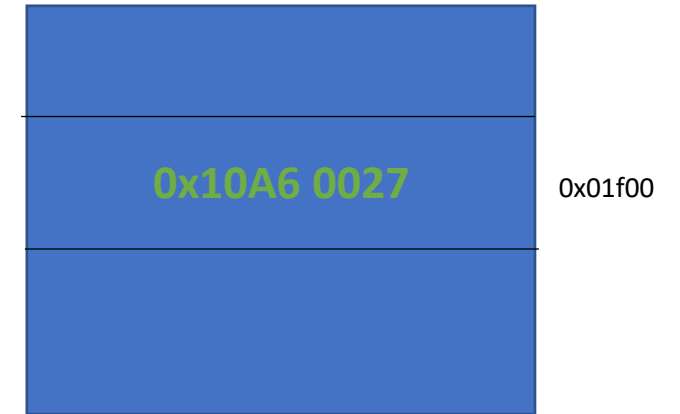
BEZ R2,100; // If R2 is zero branch to location PC+100.



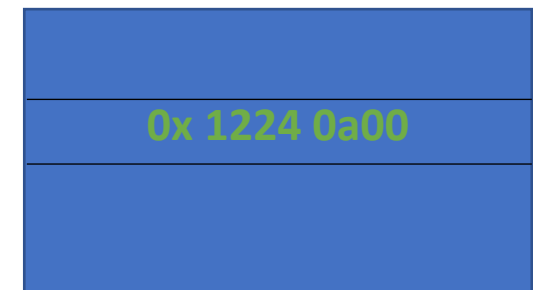
CPU

BEZ R2,100;

Program Memory

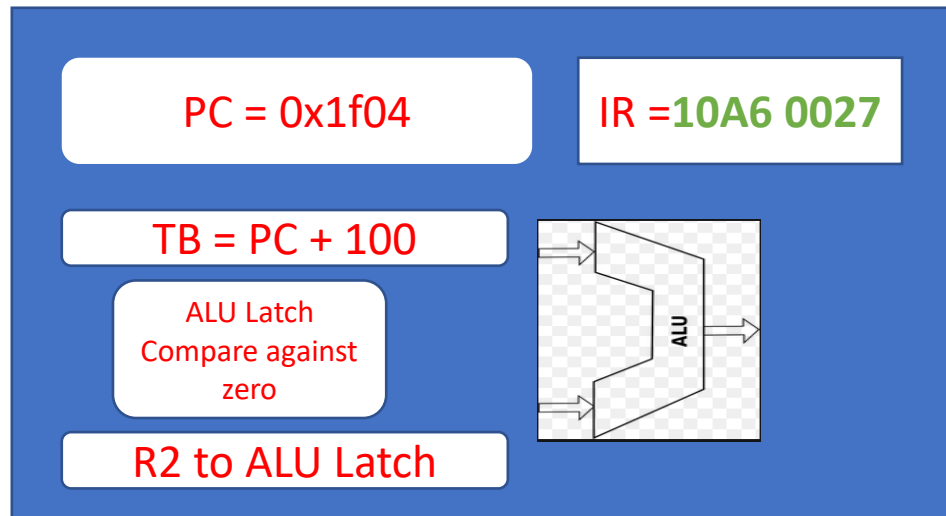


Data Memory



A 5-Stage Pipeline (Branch instruction)

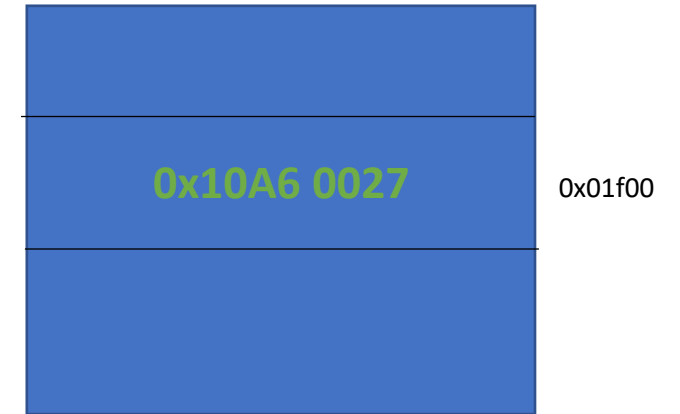
BEZ R2,100; // If R2 is zero branch to location PC+100.



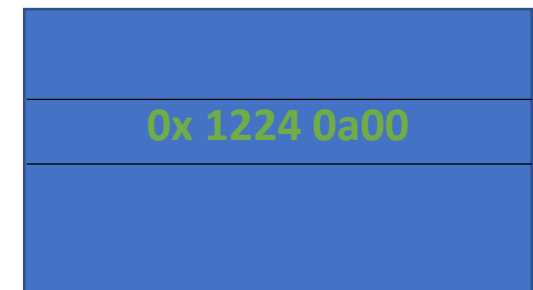
CPU

BEZ R2,100;

Program Memory



Data Memory

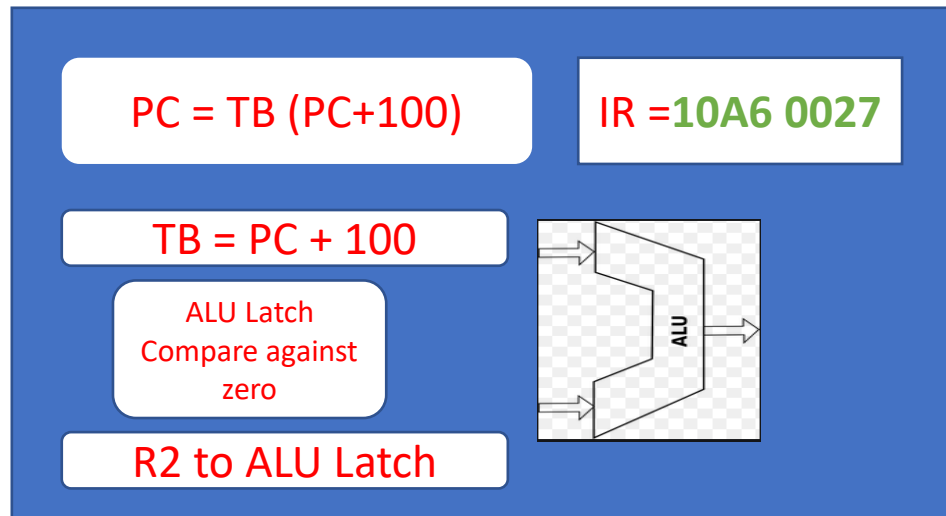


A 5-Stage Pipeline (Branch instruction)

BEZ R2,100; // If R2 is zero branch to location PC+100.



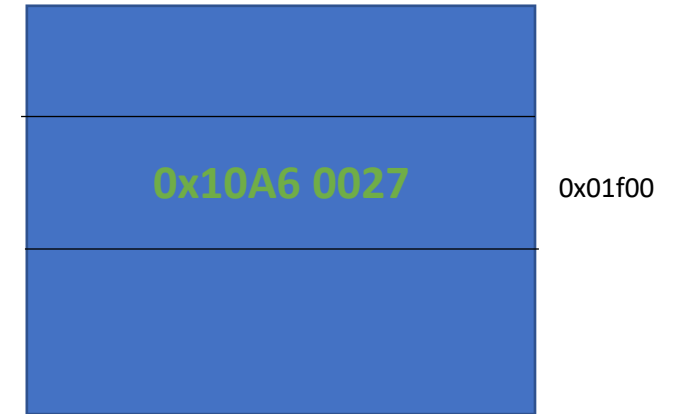
If the Branch to be taken , the PC is updated



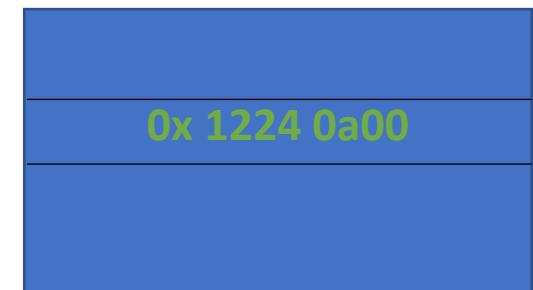
CPU

BEZ R2,100;

Program Memory



Data Memory

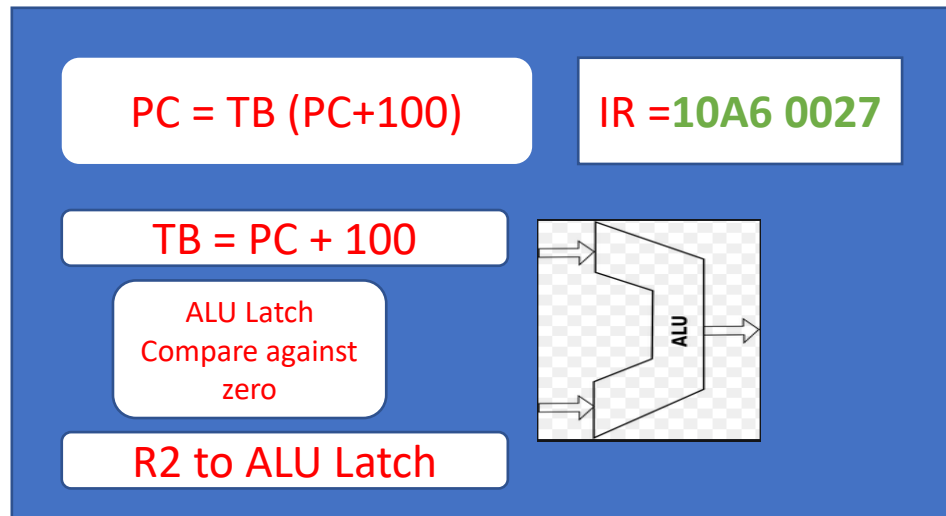


A 5-Stage Pipeline (Branch instruction)

BEZ R2,100; // If R2 is zero branch to location PC+100.



The Execute stage has no additional work to do.



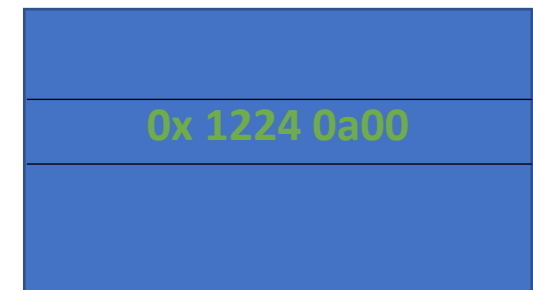
CPU

BEZ R2,100;

Program Memory



Data Memory

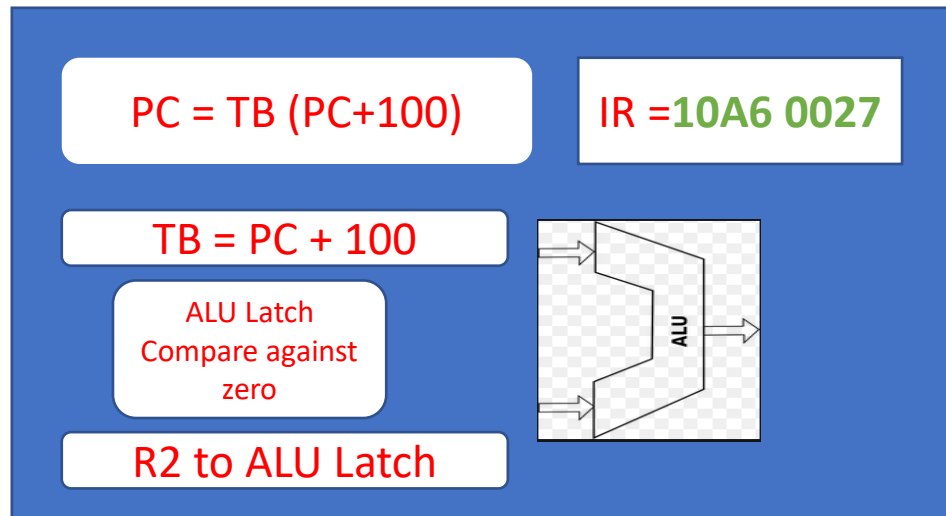


A 5-Stage Pipeline (Branch instruction)

BEZ R2,100; // If R2 is zero branch to location PC+100.



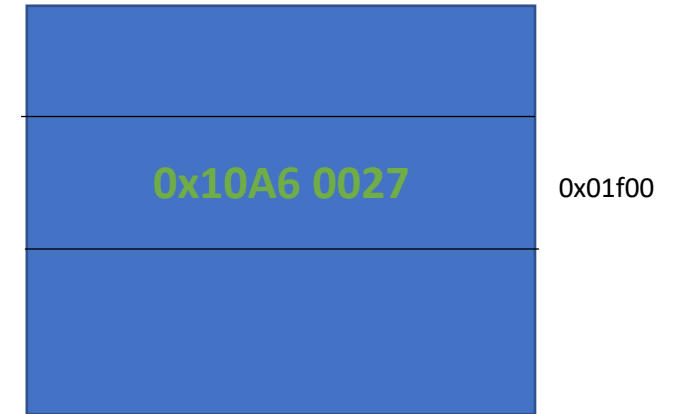
The Memory stage has no additional work to do.



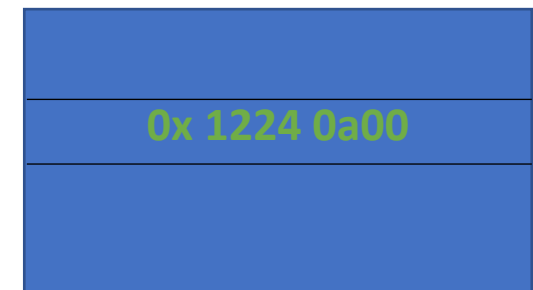
CPU

BEZ R2,100;

Program Memory



Data Memory

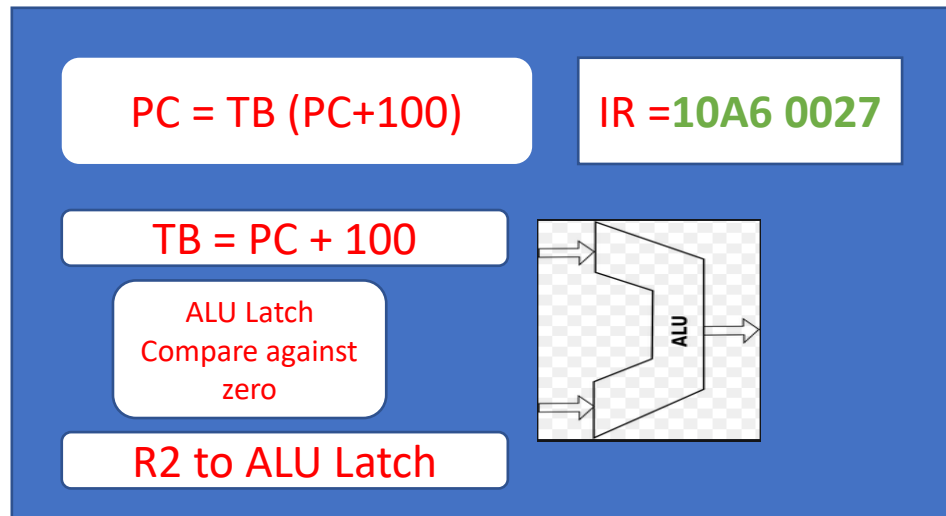


A 5-Stage Pipeline (Branch instruction)

BEZ R2,100; // If R2 is zero branch to location PC+100.



The Store stage has no additional work to do.



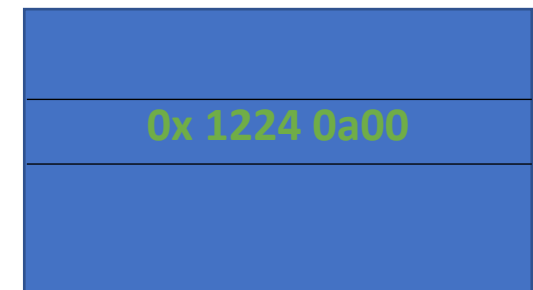
CPU

BEZ R2,100;

Program Memory



Data Memory



RISC: Five-Stage Pipeline

Instruction number	Clock number								
	1	2	3	4	5	6	7	8	9
Instruction i	IF	ID	EX	MEM	WB				
Instruction $i + 1$		IF	ID	EX	MEM	WB			
Instruction $i + 2$			IF	ID	EX	MEM	WB		
Instruction $i + 3$				IF	ID	EX	MEM	WB	
Instruction $i + 4$					IF	ID	EX	MEM	WB

Simply start a new instruction
on each clock cycle;
Speedup = 5.

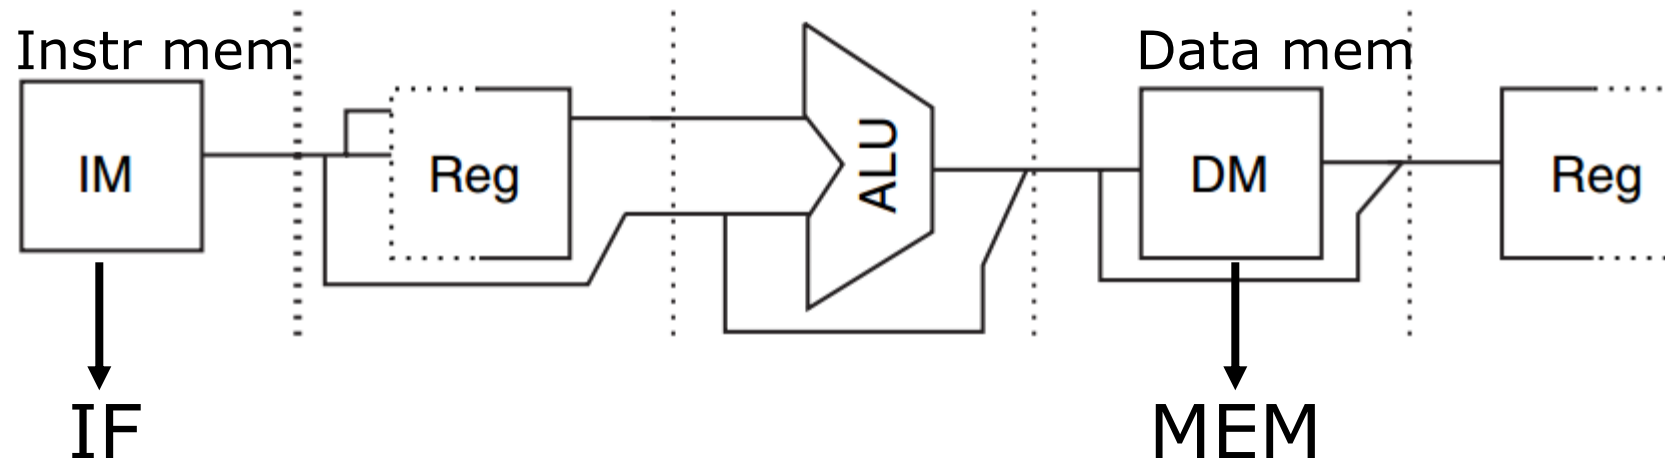
RISC: Five-Stage Pipeline

- To achieve the pipelining, we need separate instruction memory and data memory.
- Because both IF and MEM require memory access, if there is only one single memory, IF and MEM will have memory conflicts during the pipelining process.

RISC: Five-Stage Pipeline

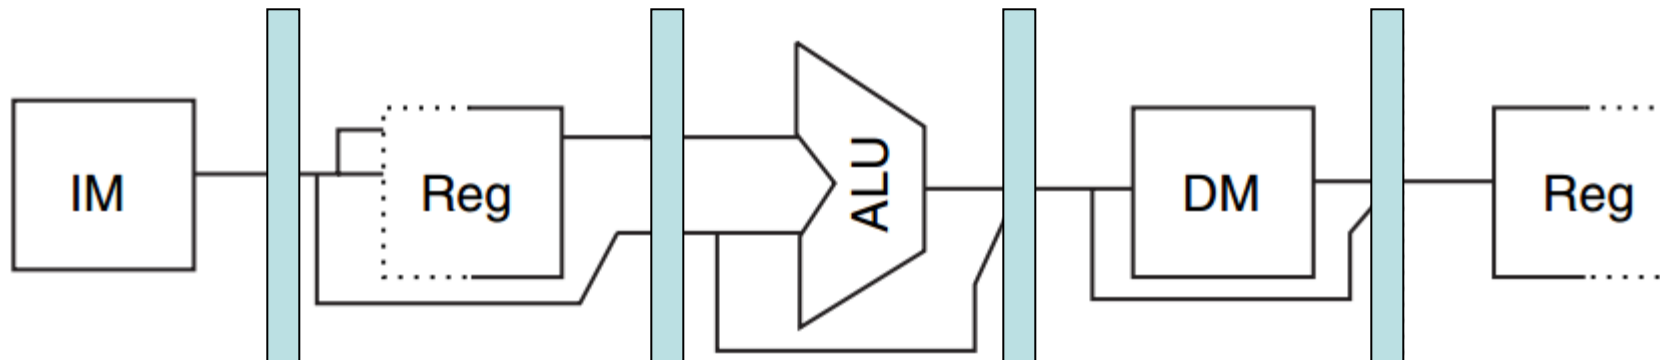
separate instruction and data mems to eliminate conflicts for a single memory between instruction fetch and data memory access.

4
MEM
EX
ID
IF



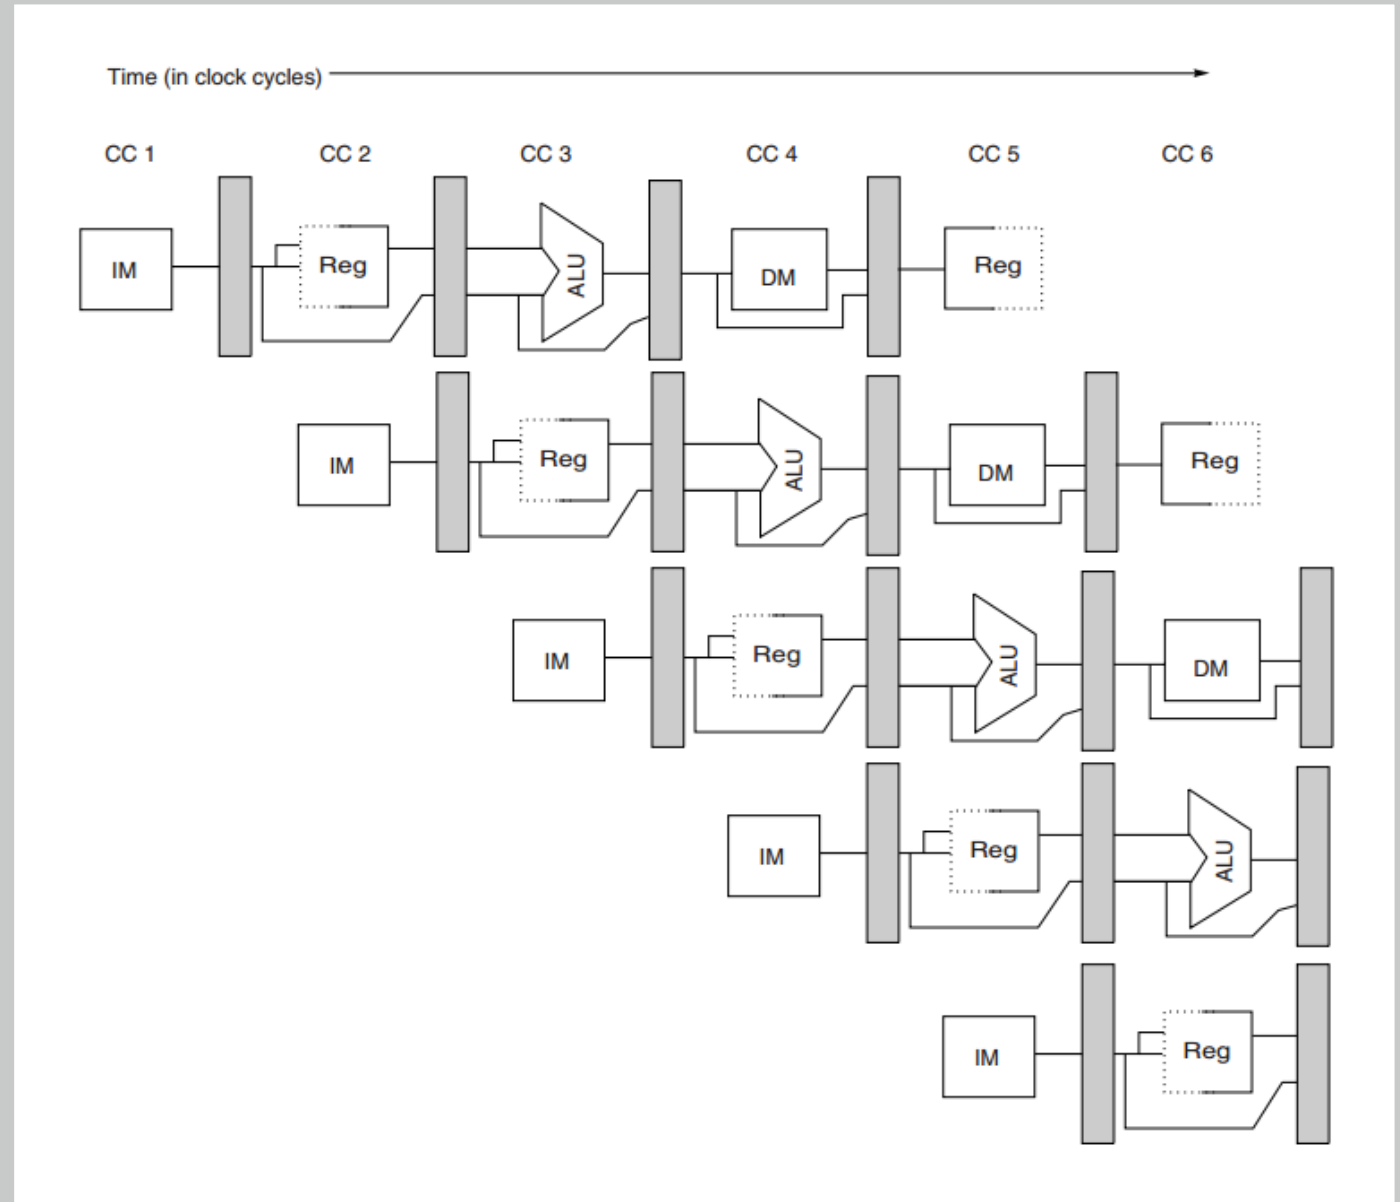
RISC: Five-Stage Pipeline

- How it works
introduce pipeline registers between successive stages;
pipeline registers store the results of a stage and use them as the input of the next stage.

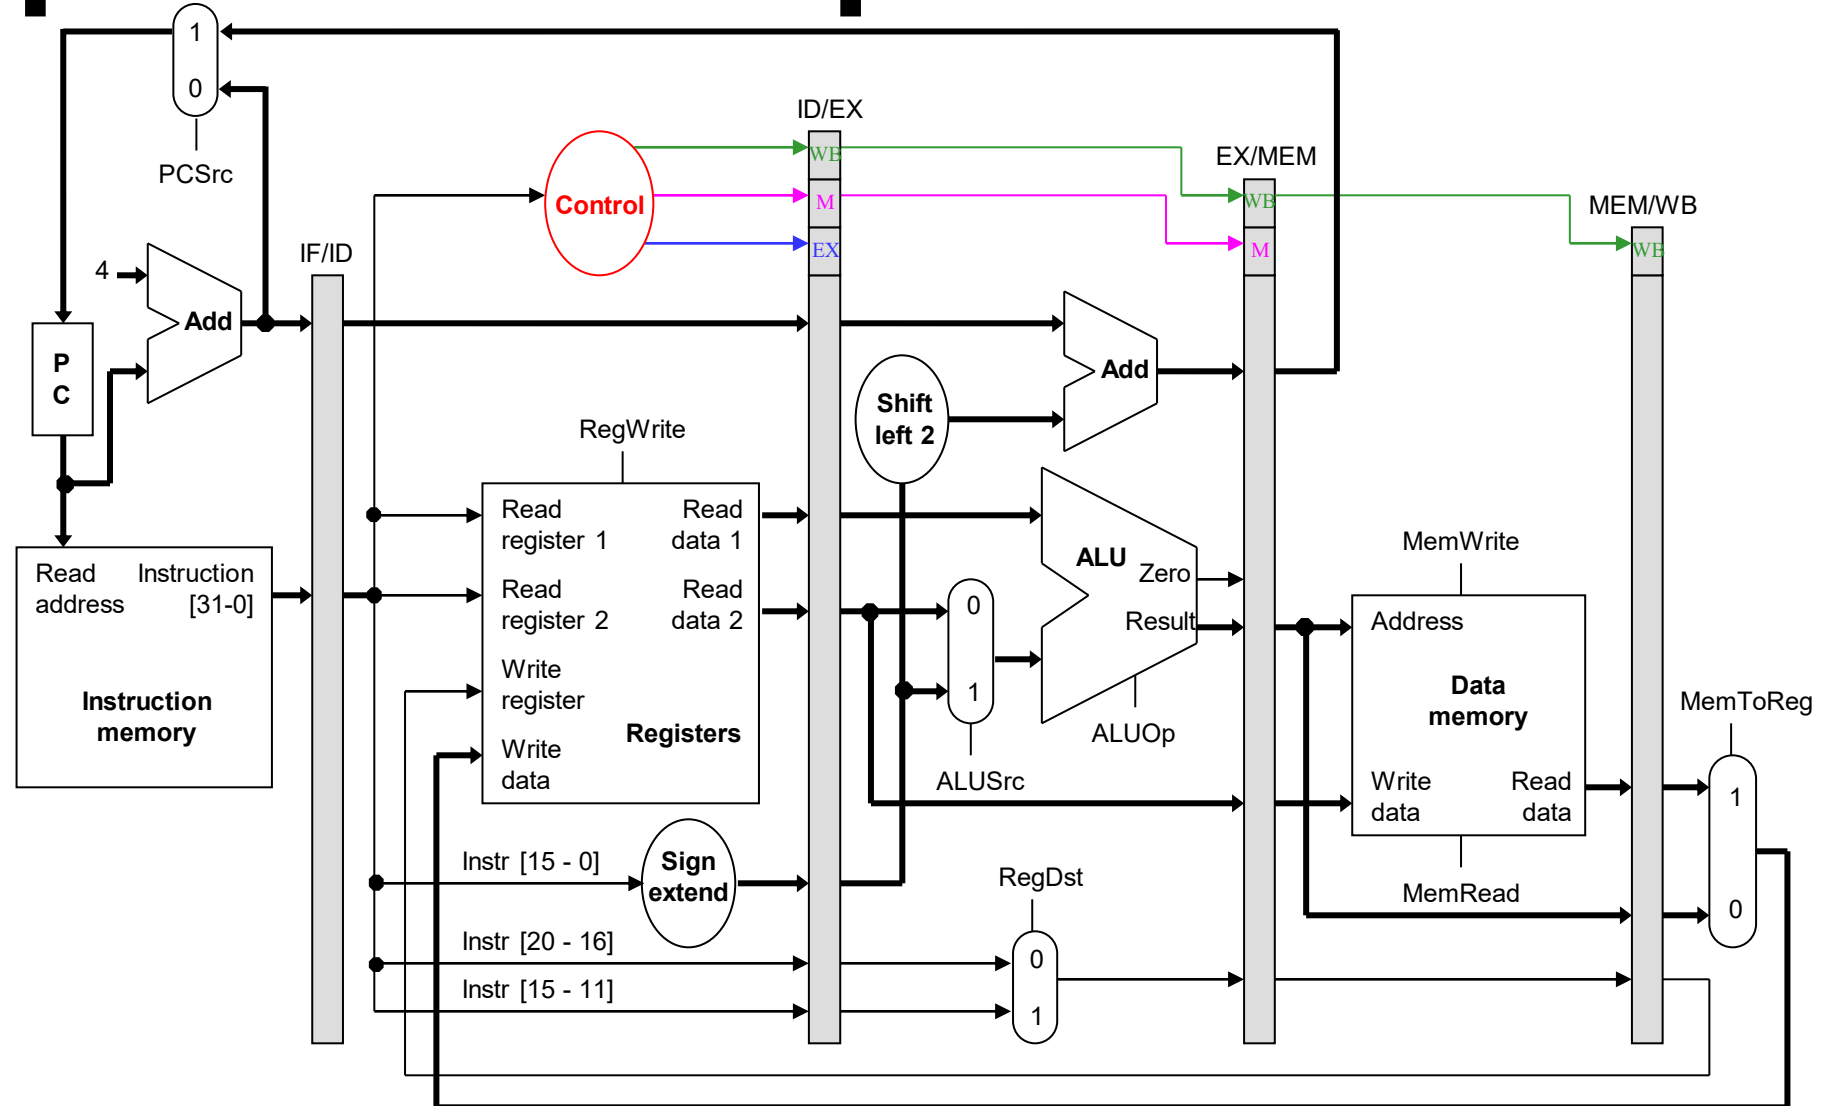


RISC: Five-Stage Pipeline

- high level paradigm of RISC pipelining



Pipelined datapath and control



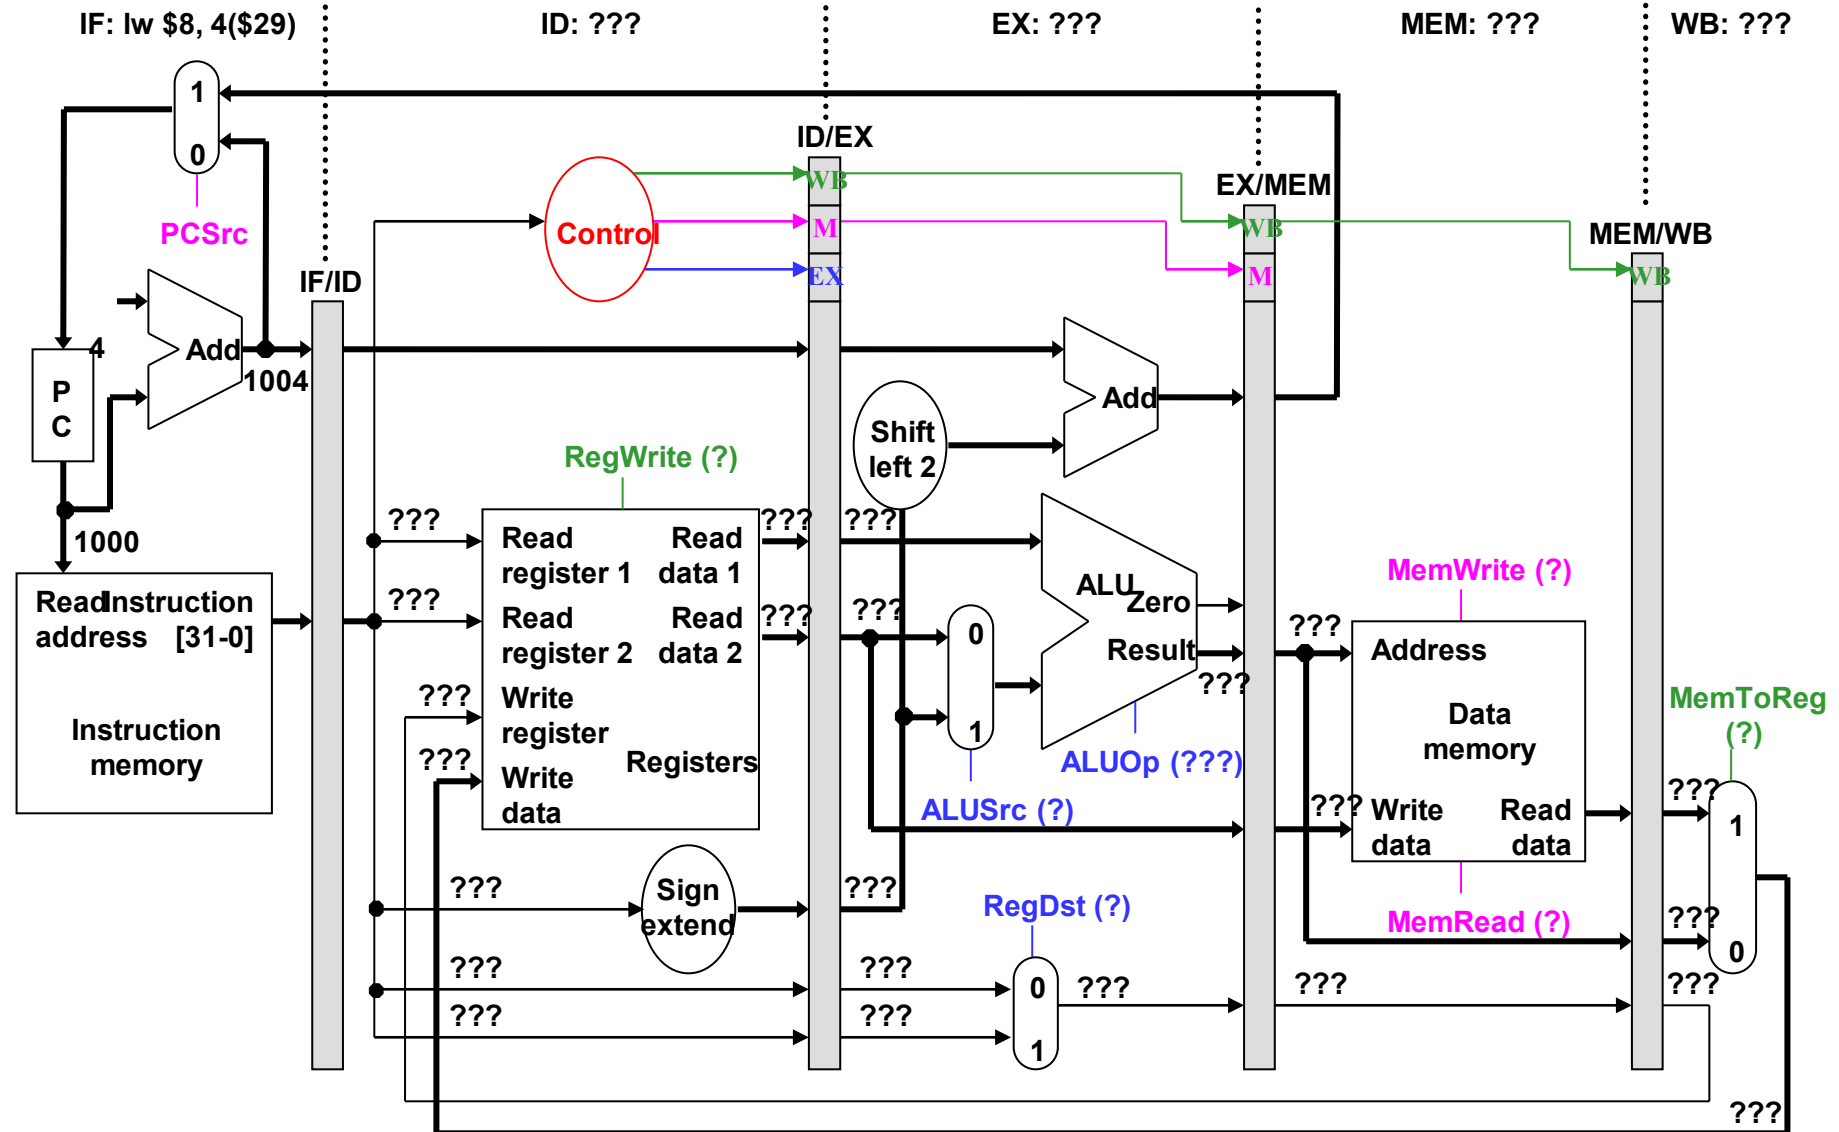
An example execution sequence

- Here's a sample sequence of instructions to execute

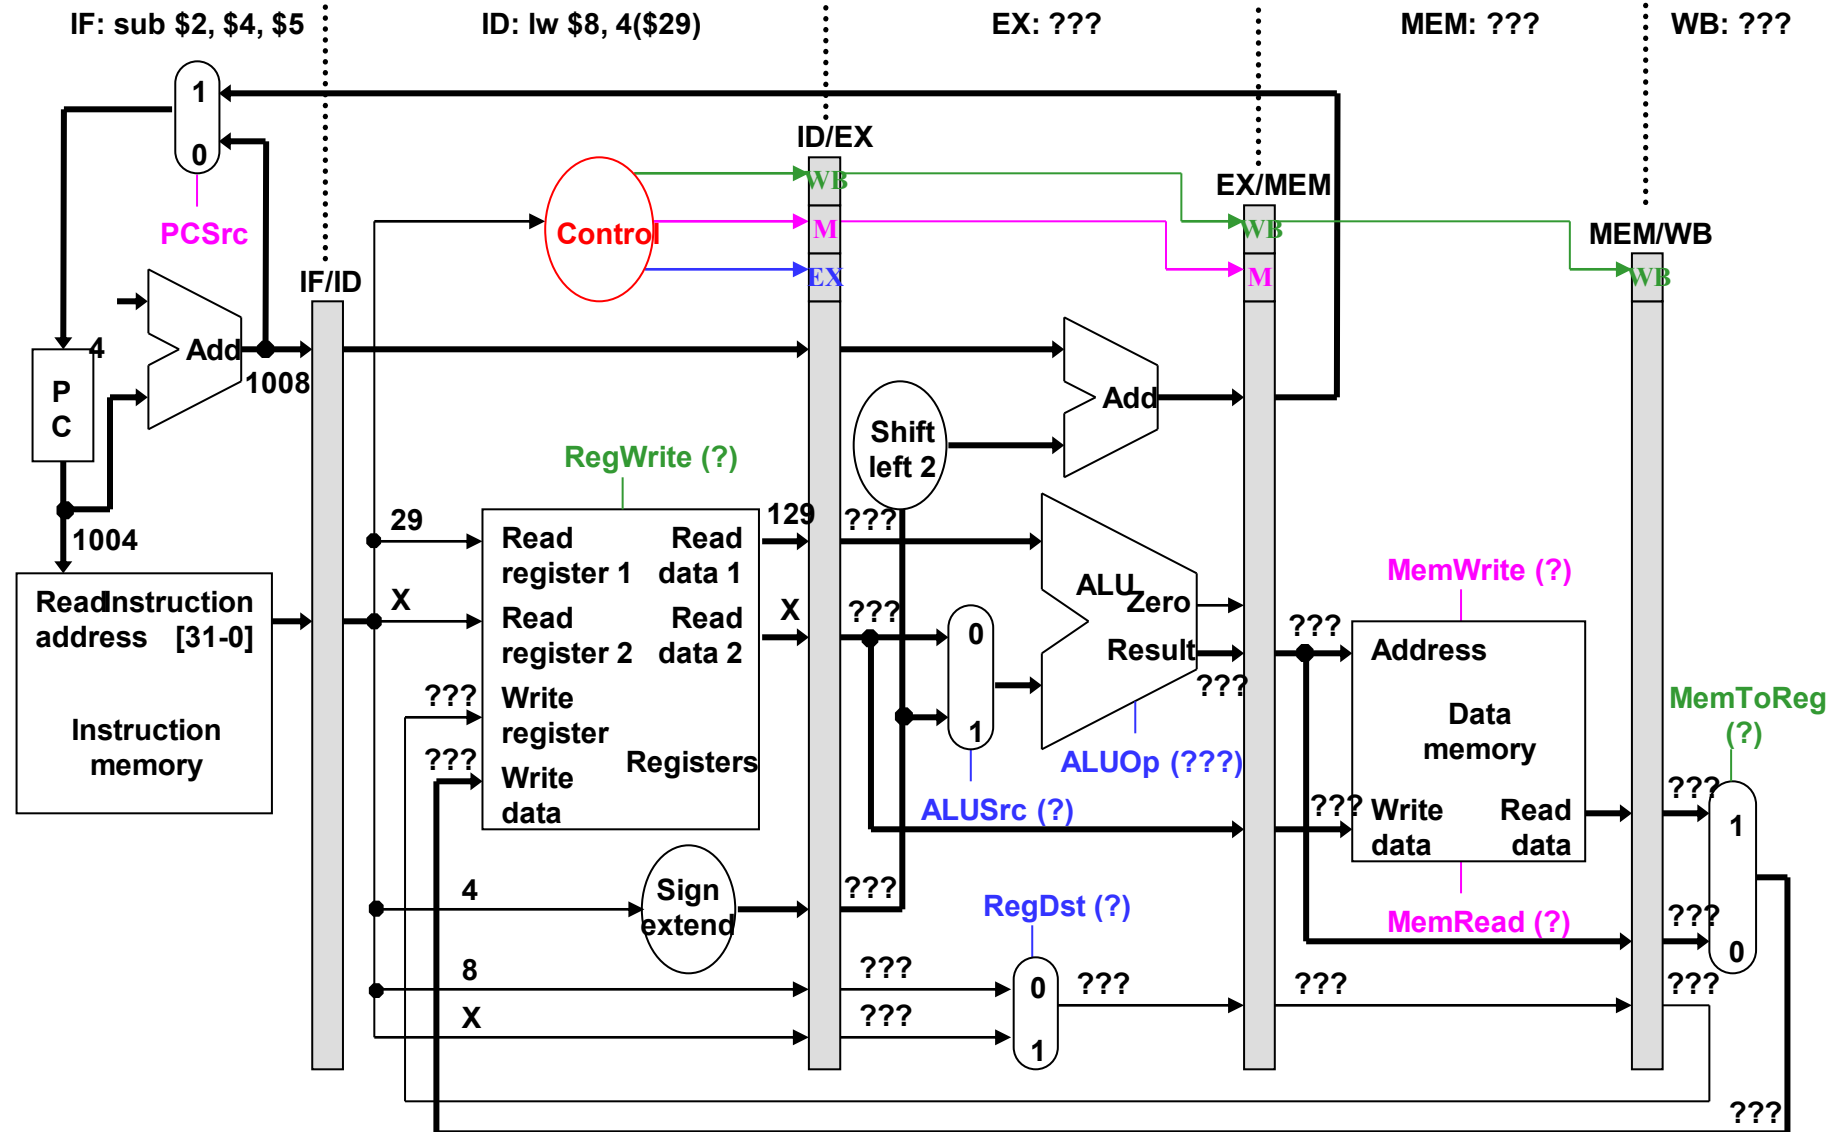
```
1000: lw  $8, 4($29)
1004: sub $2, $4, $5
1008: and $9, $10, $11
1012: or  $16, $17, $18
1016: add $13, $14, $0
```

- We'll make some assumptions, just so we can show actual data values:
 - Each register contains its number plus 100. For instance, register \$8 contains 108, register \$29 contains 129, etc.
 - Every data memory location contains 99
- Our pipeline diagrams will follow some conventions:
 - An **X** indicates values that aren't important, like the constant field of an R-type instruction
 - Question marks **???** indicate values we don't know, usually resulting from instructions coming before and after the ones in our example

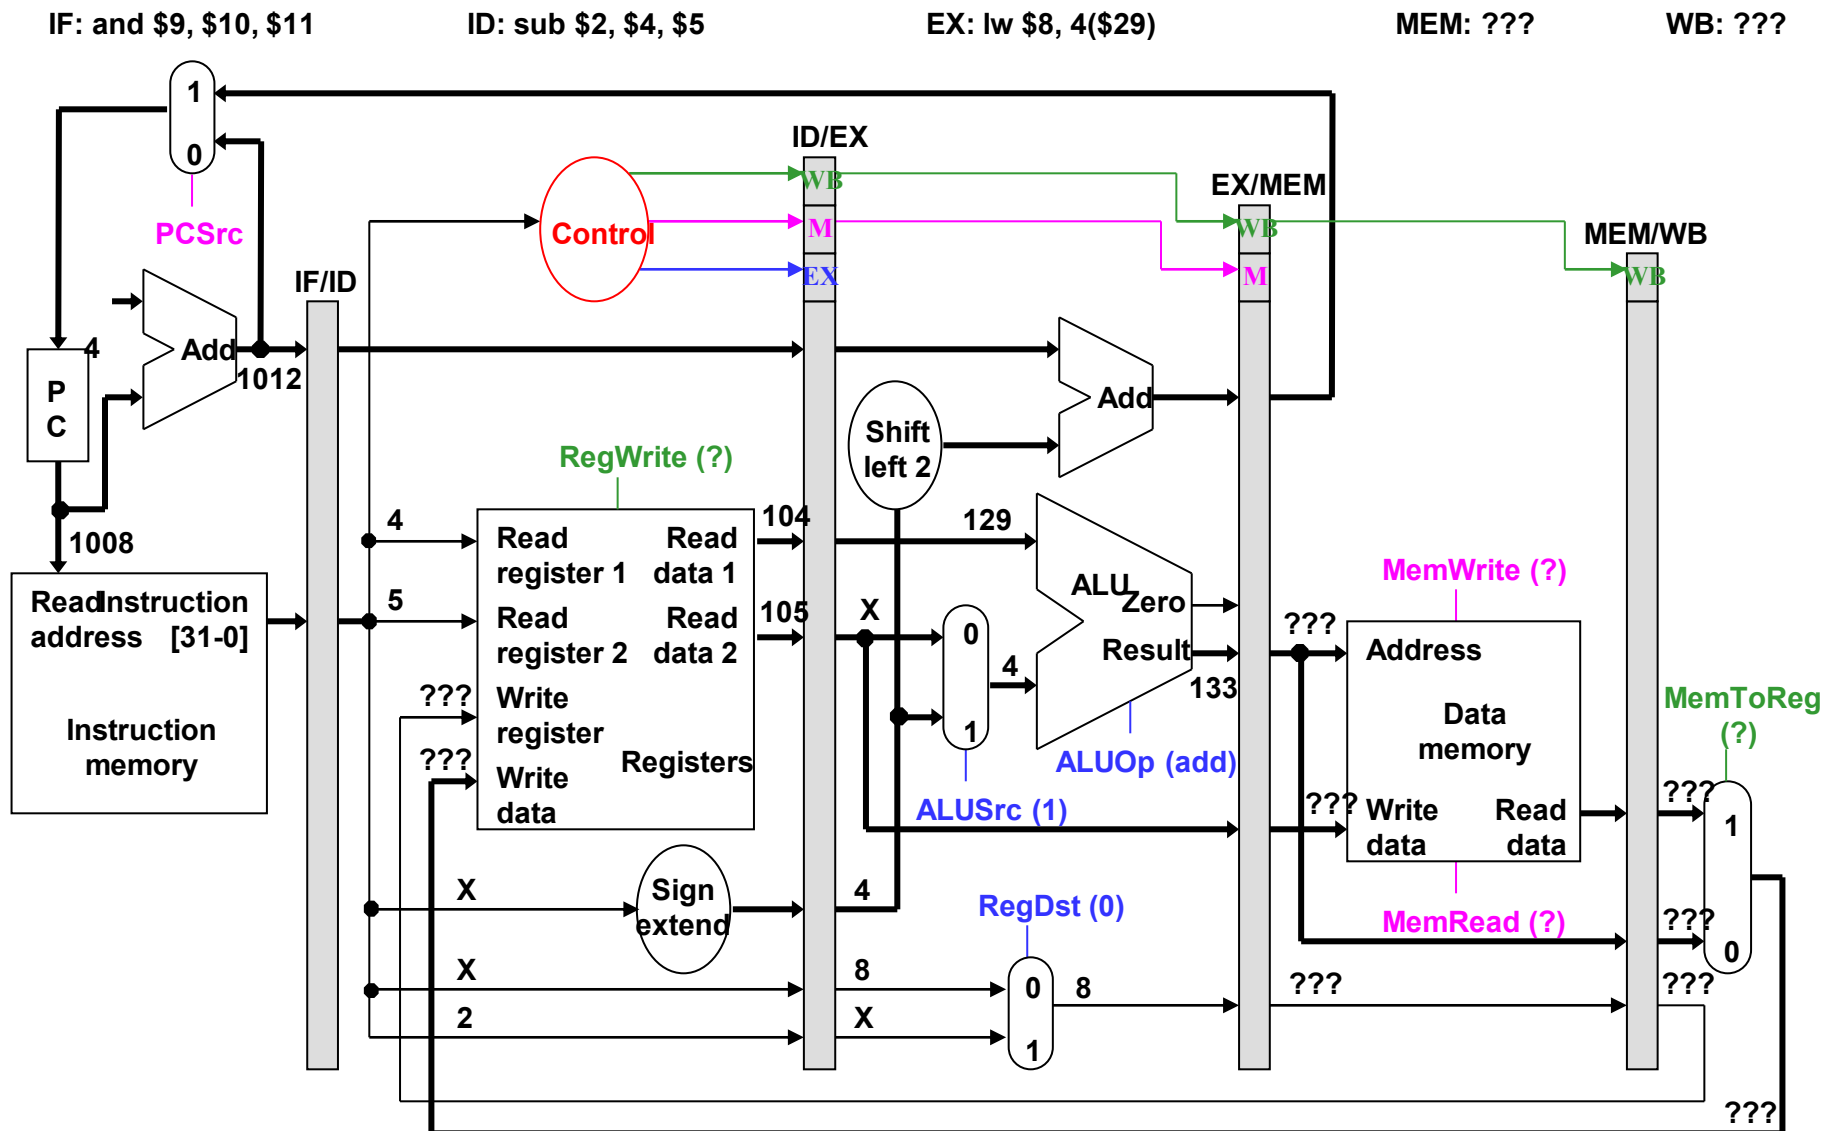
Cycle 1 (filling)



Cycle 2



Cycle 3



Cycle 4

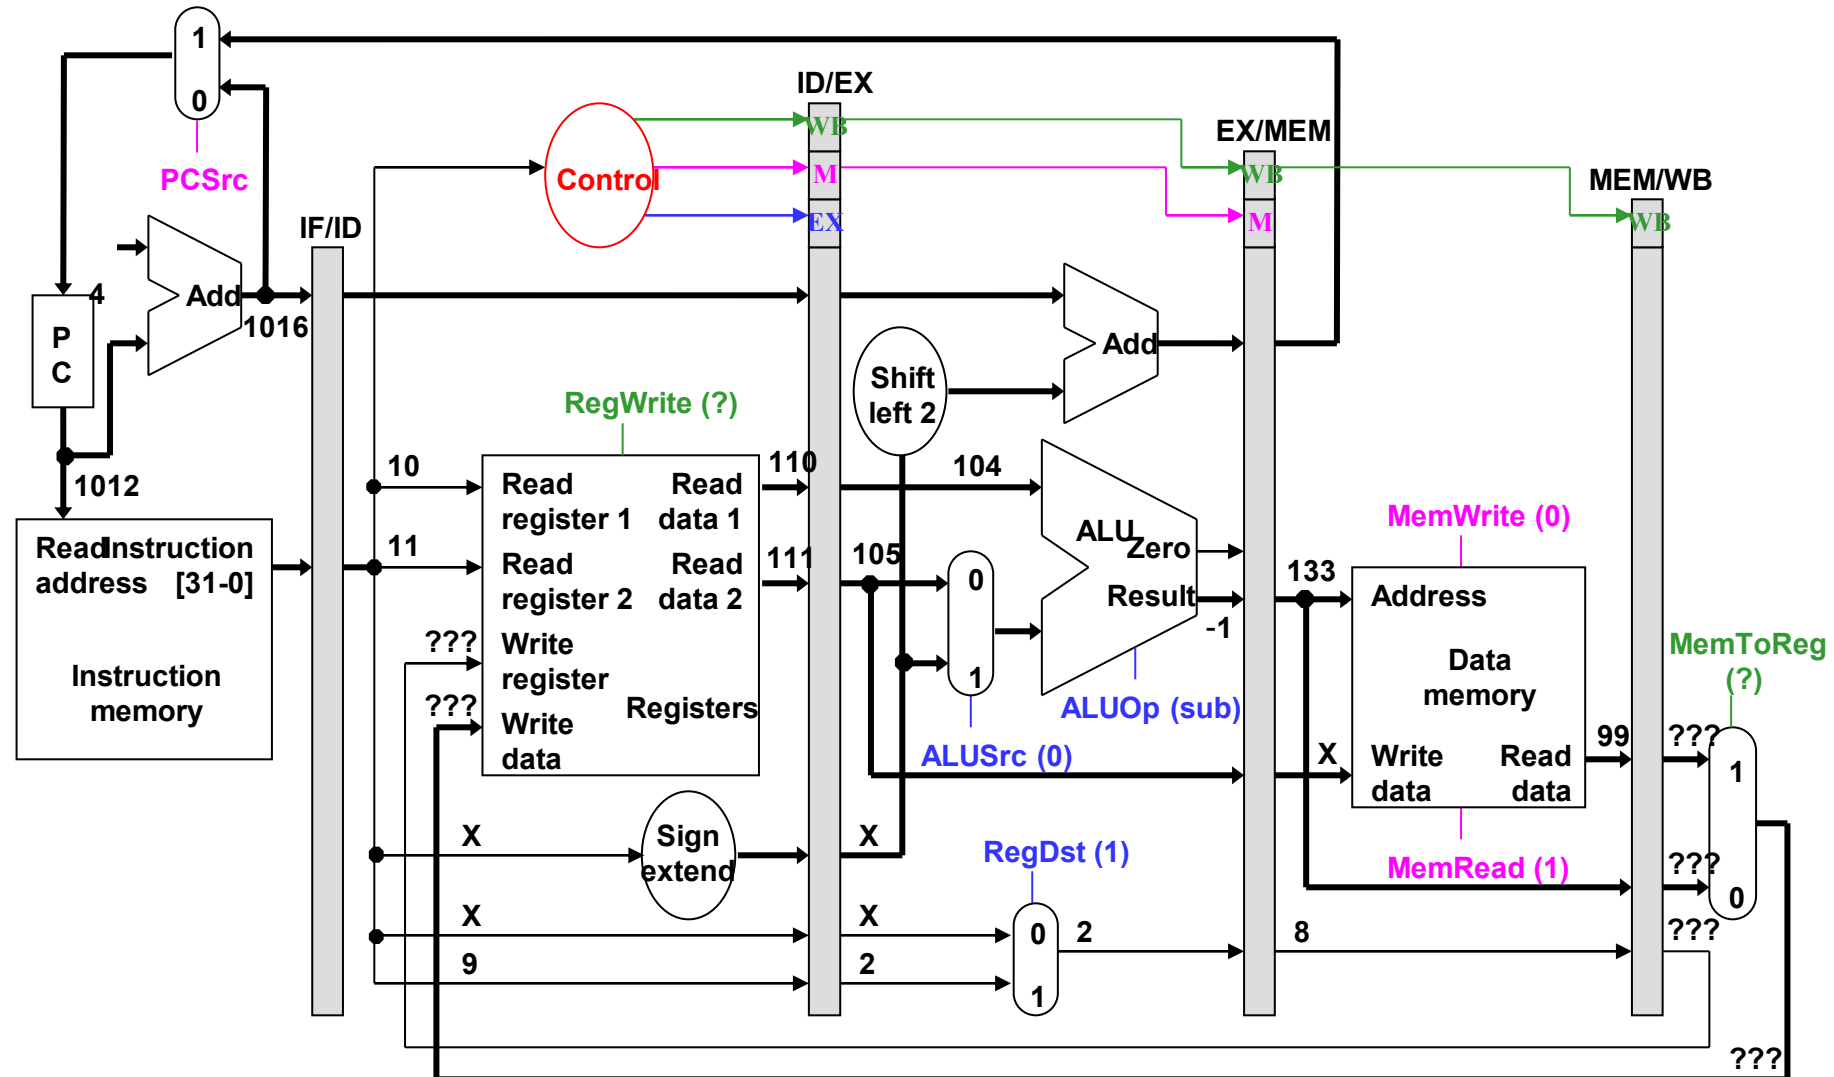
IF: or \$16, \$17, \$18

ID: and \$9, \$10, \$11

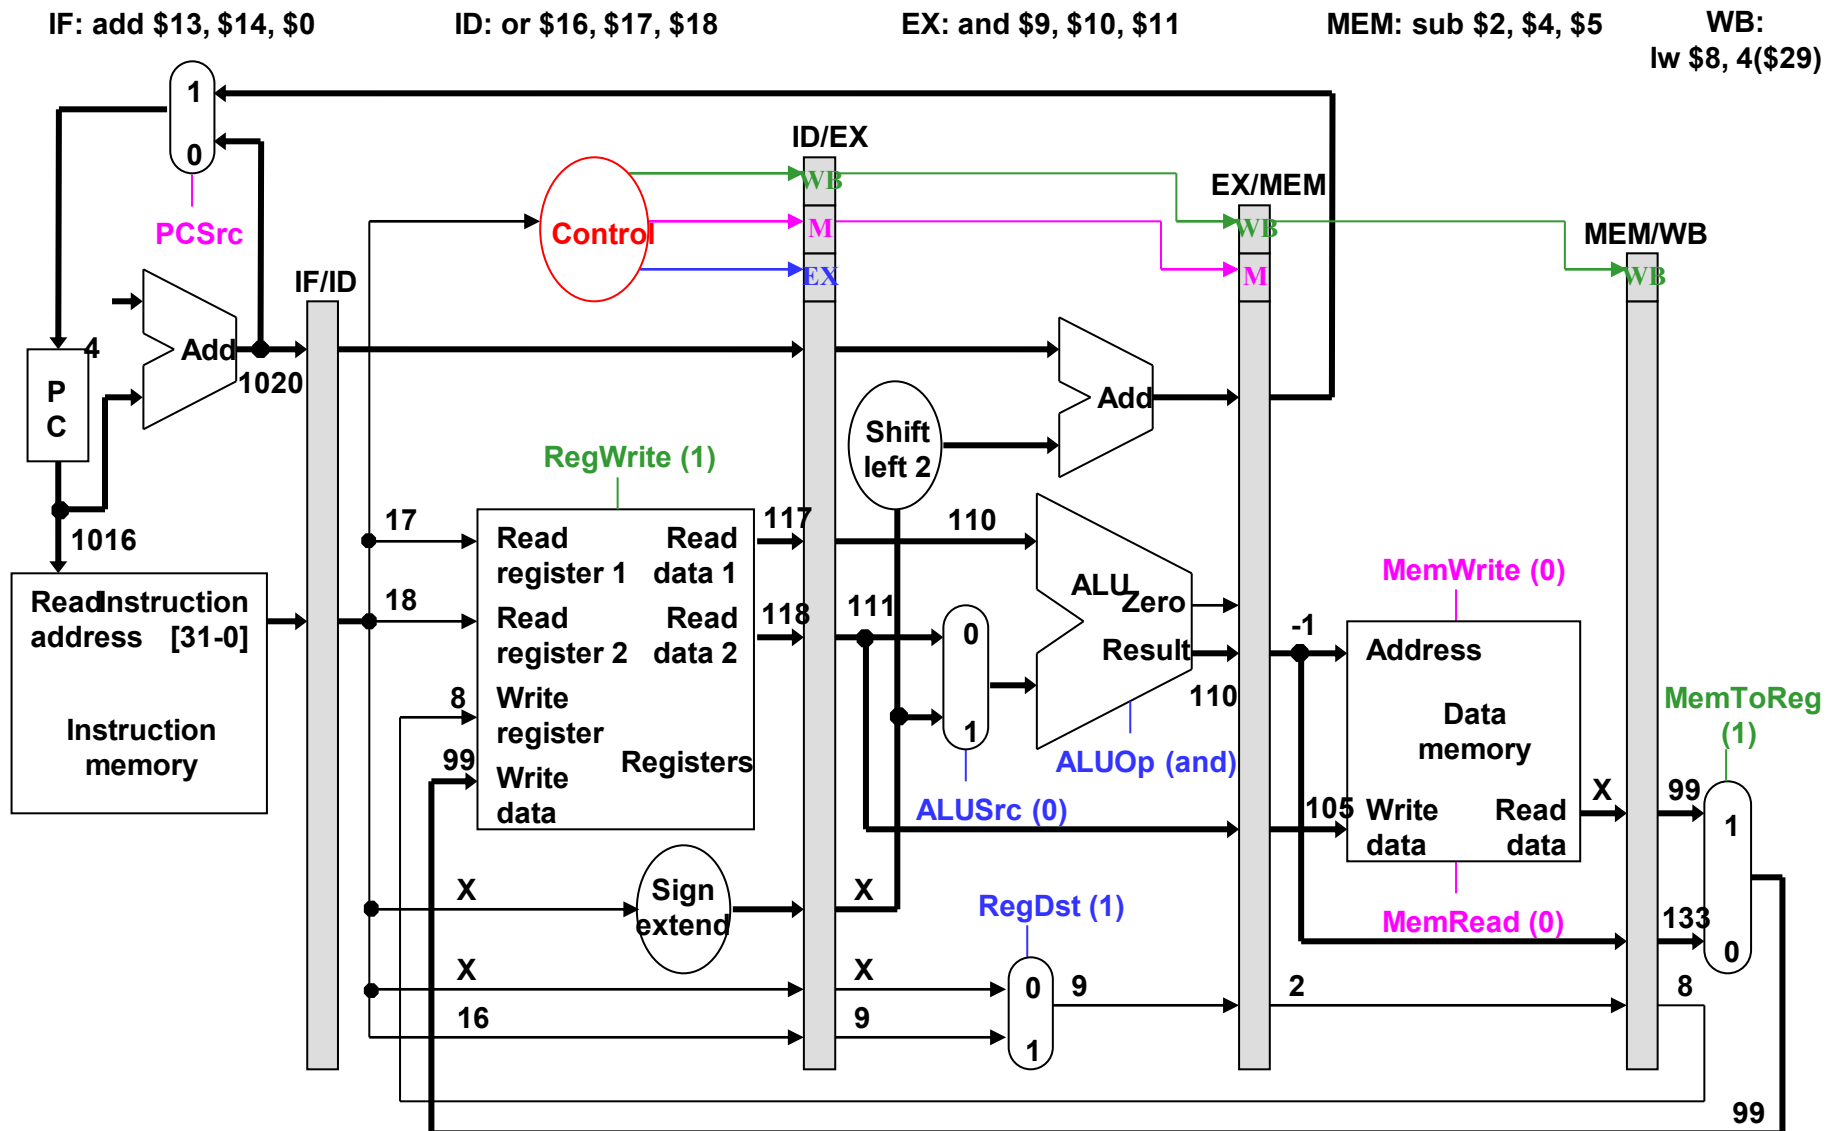
EX: sub \$2, \$4, \$5

MEM: lw \$8, 4(\$29)

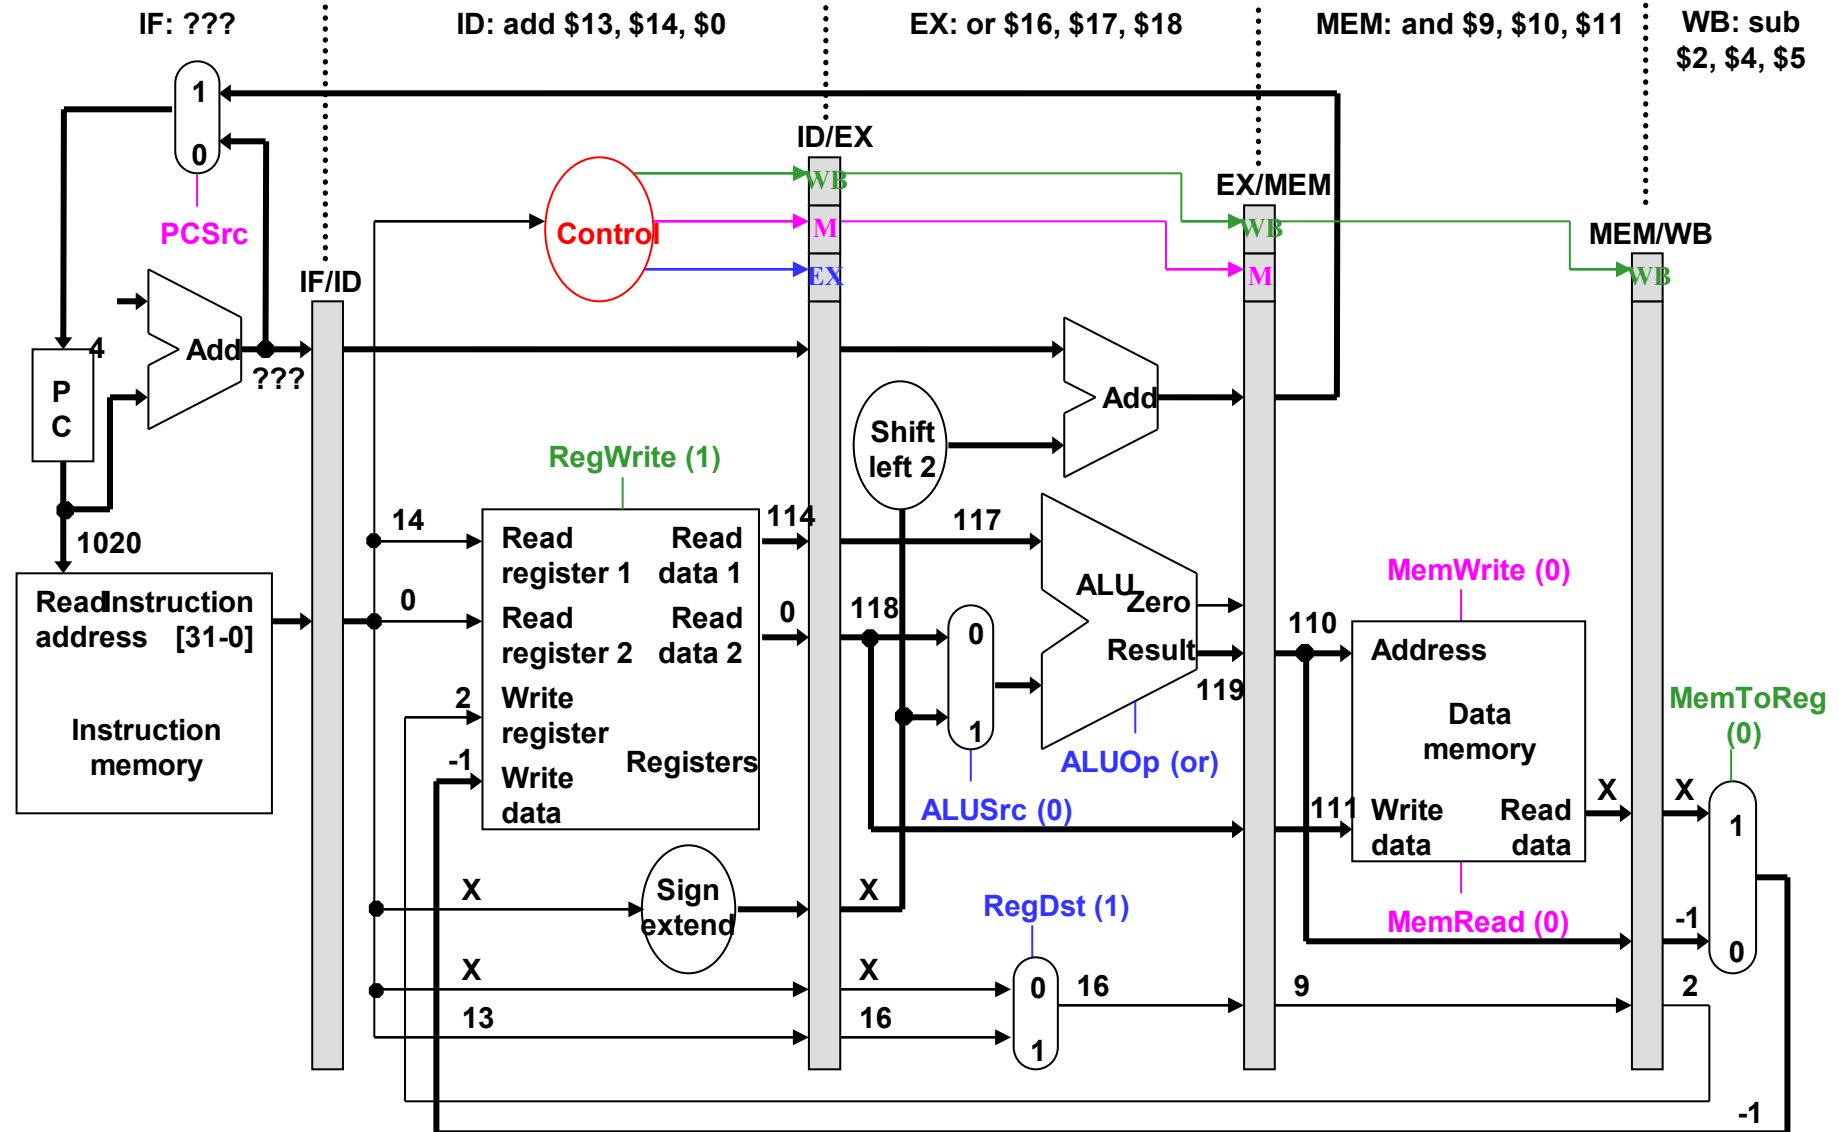
WB: ???



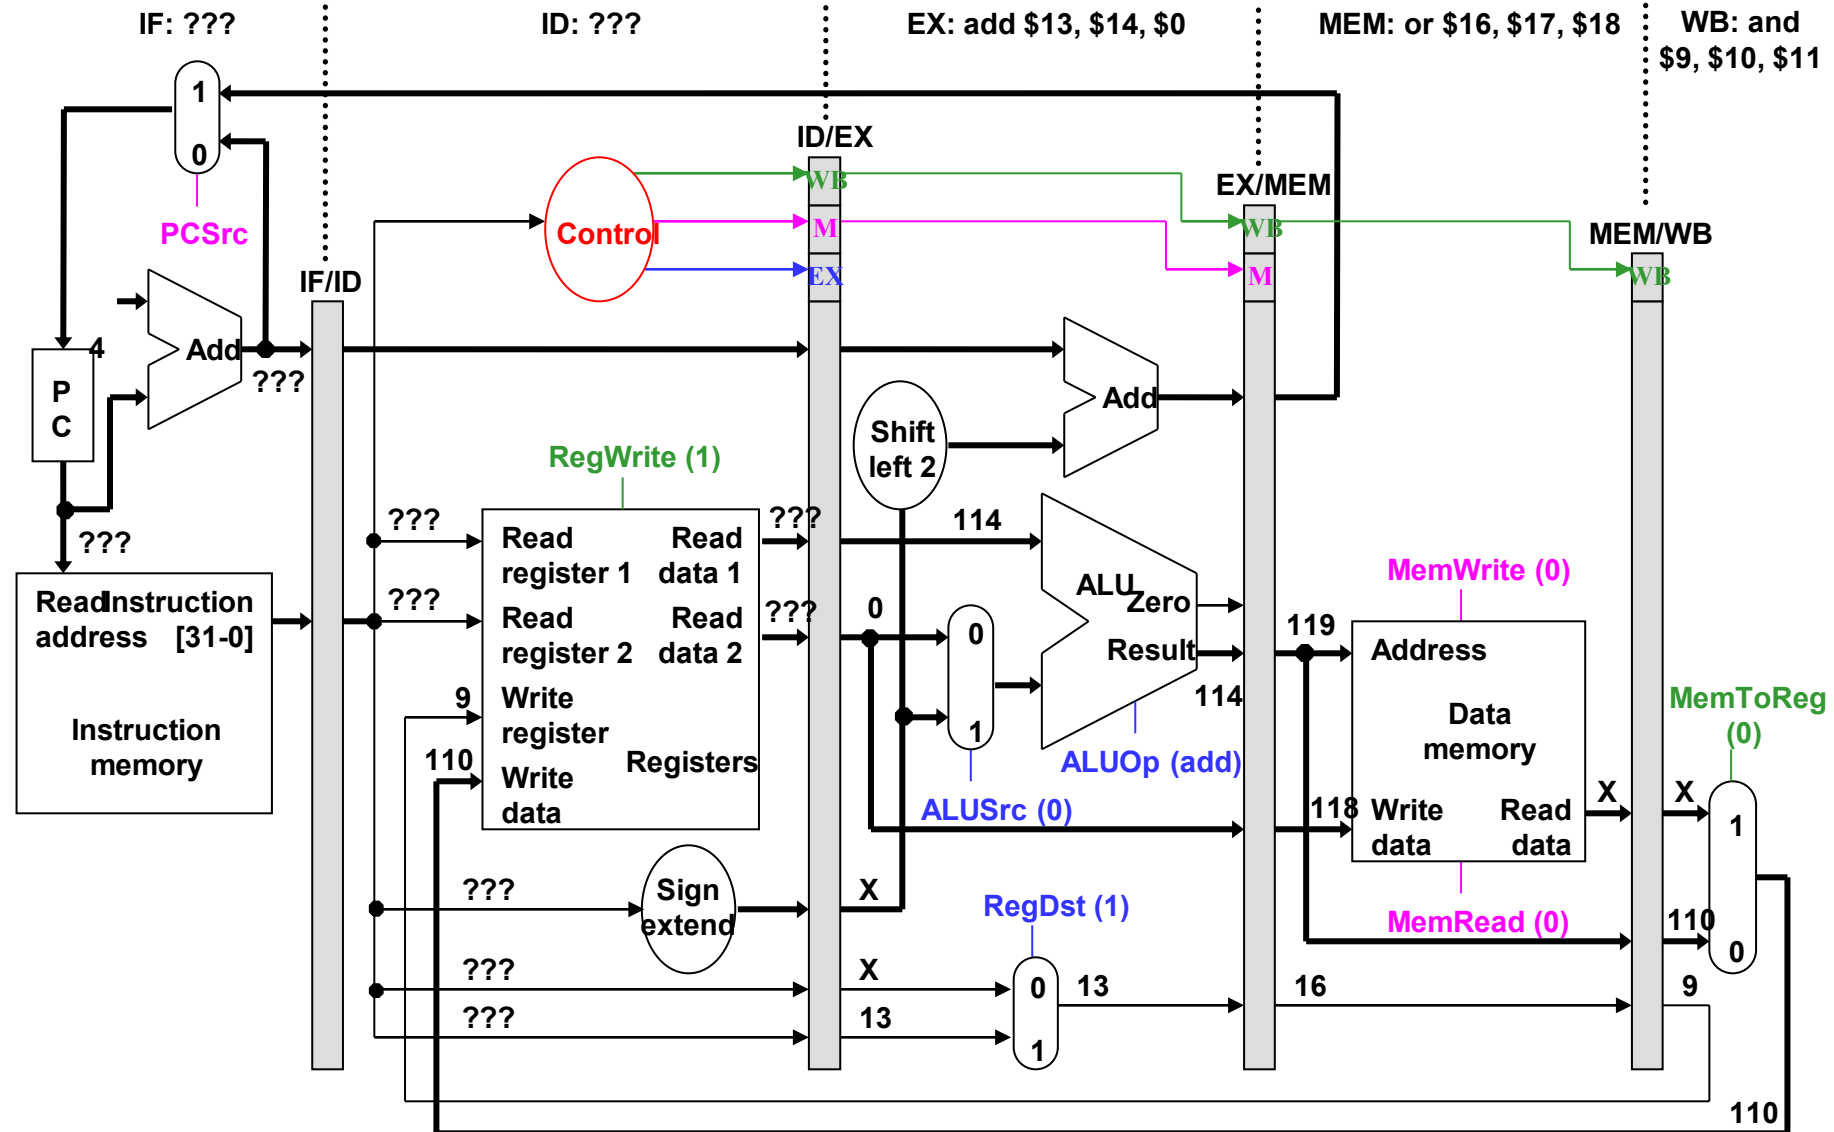
Cycle 5 (full)



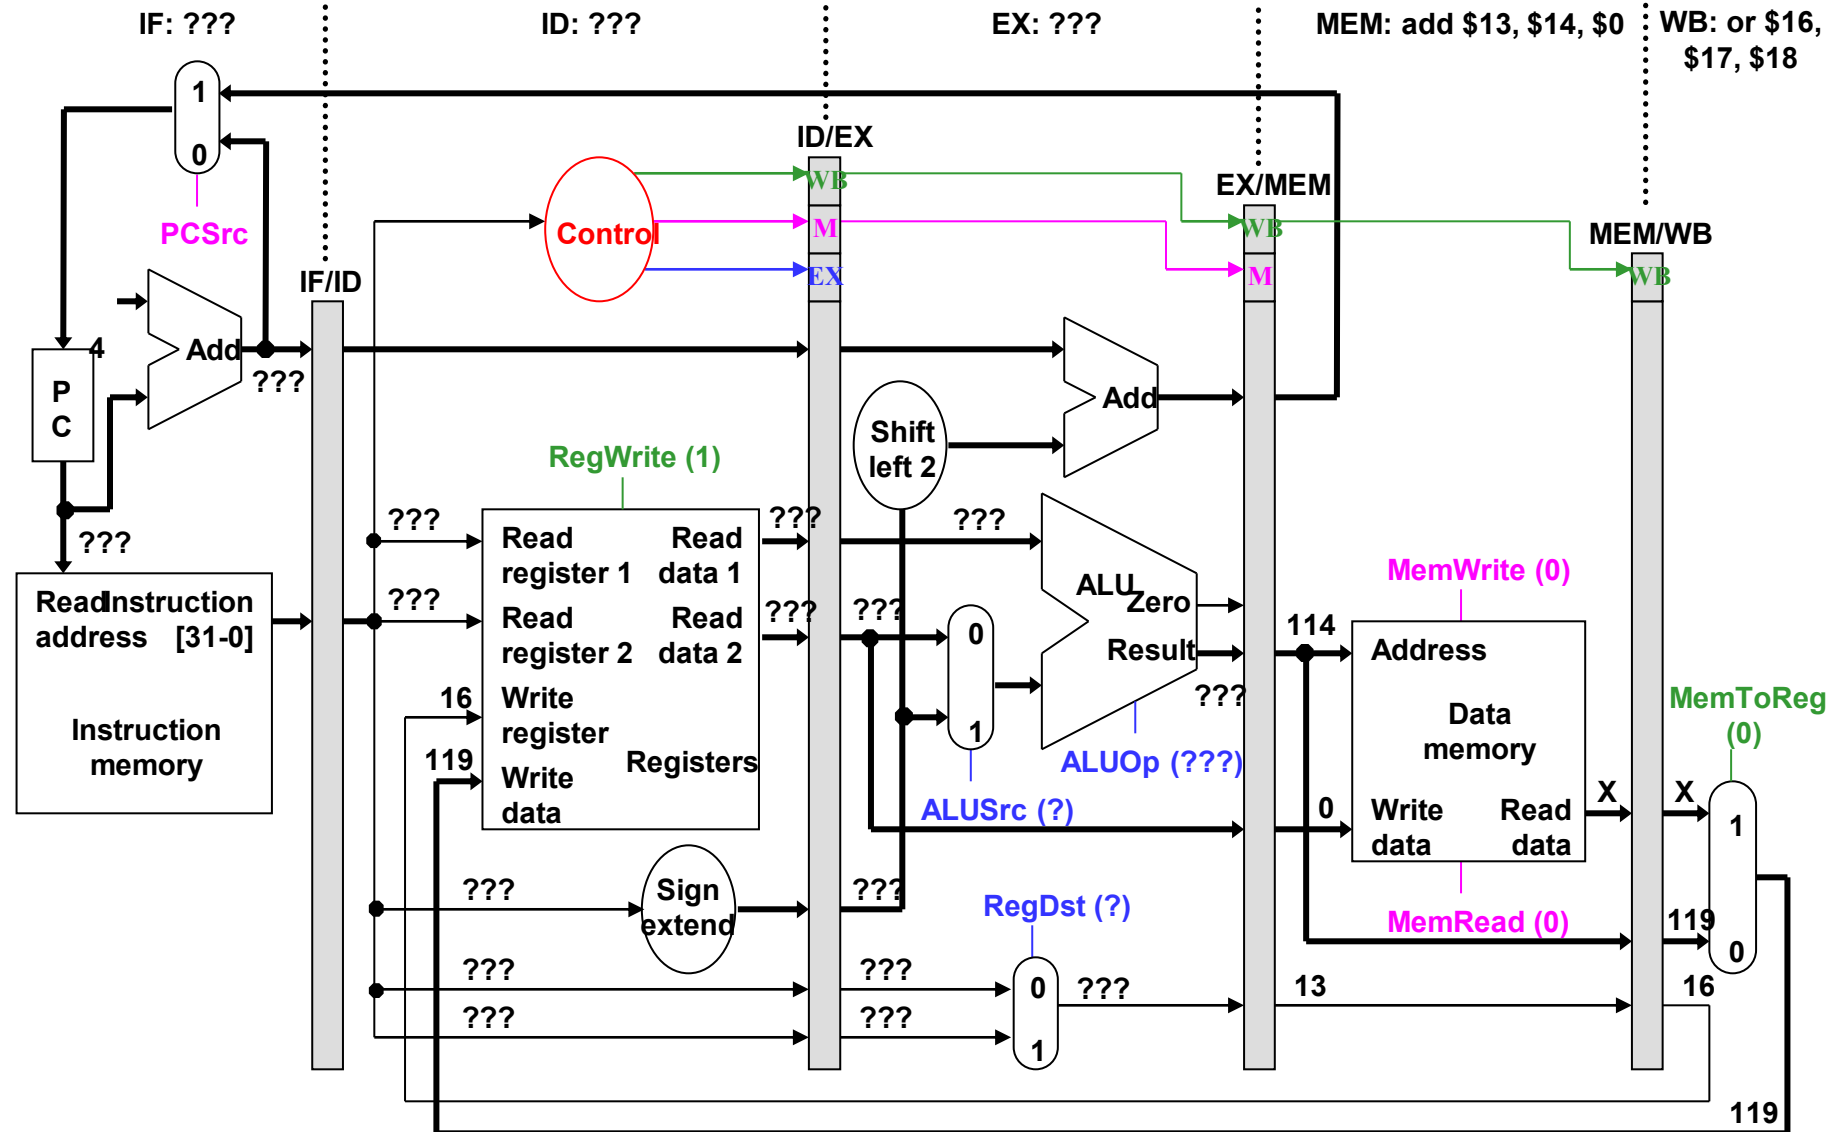
Cycle 6 (emptying)



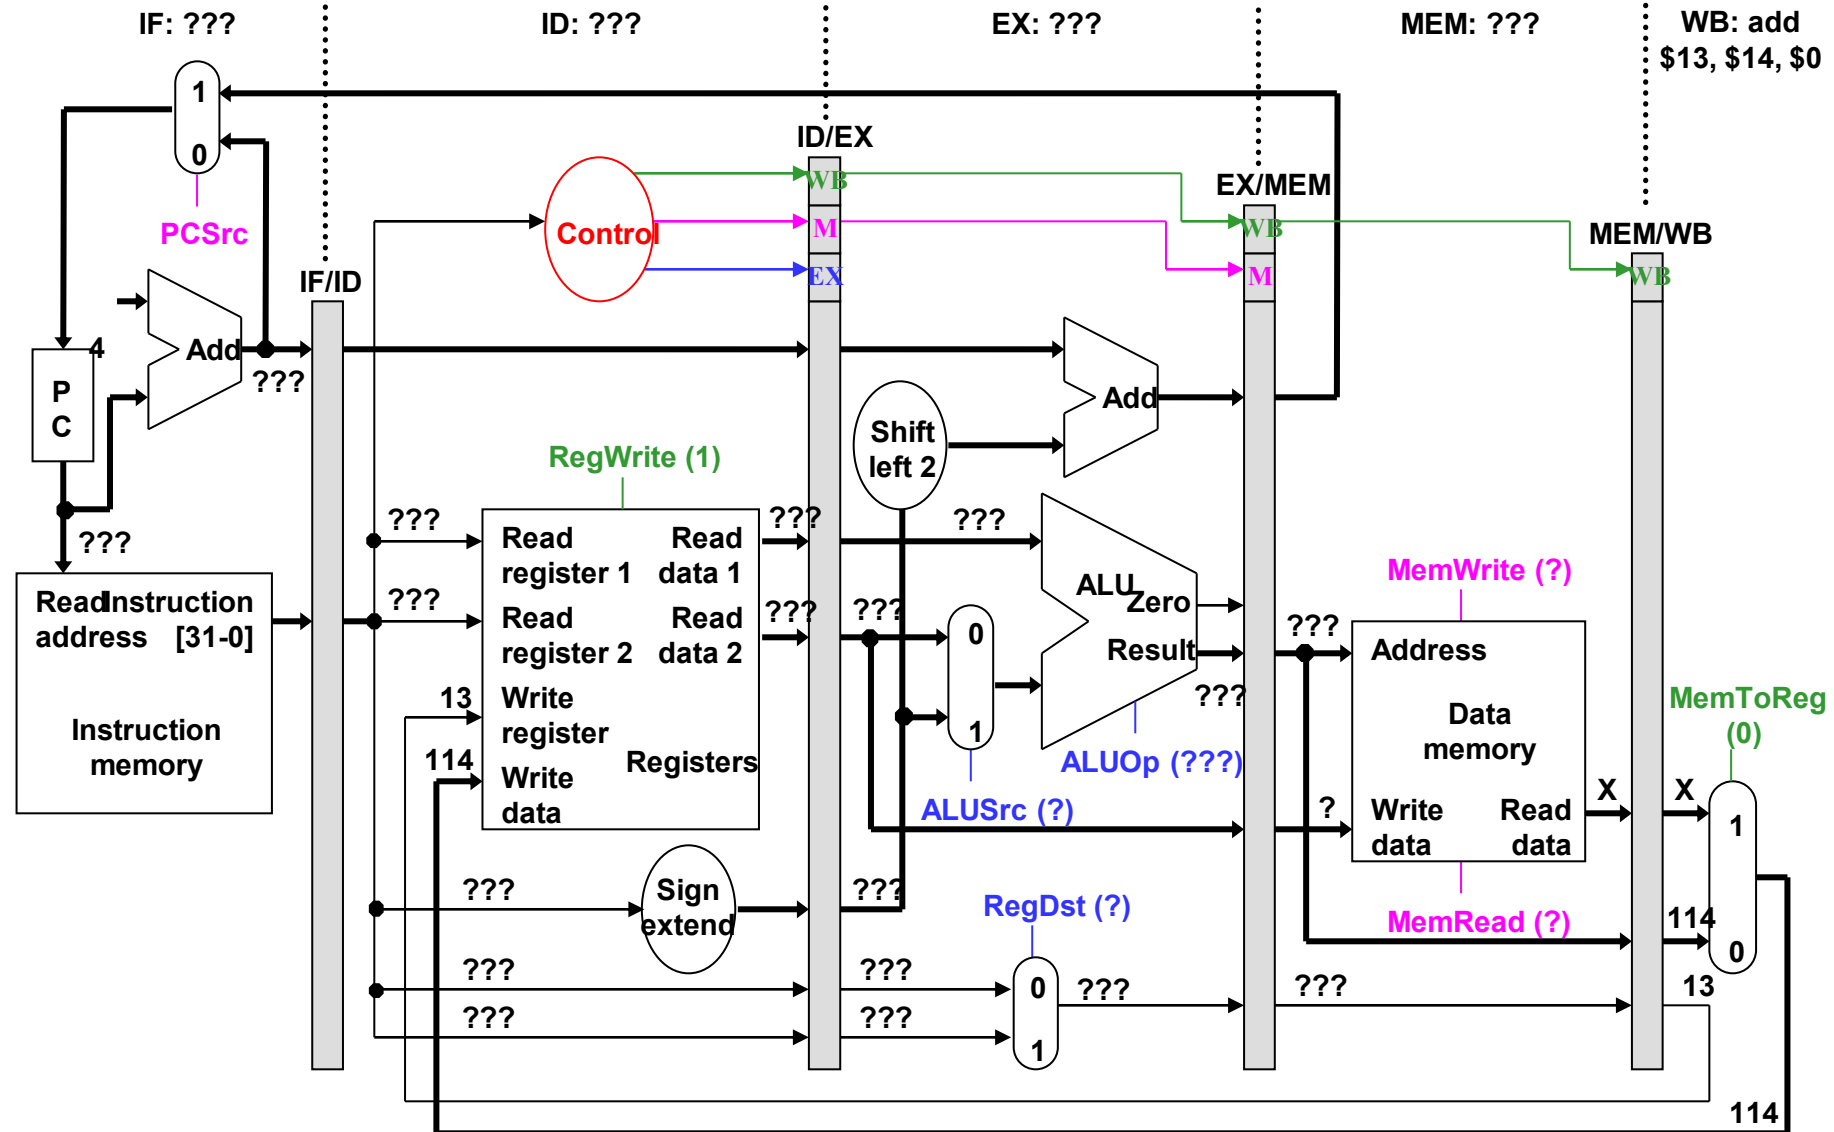
Cycle 7



Cycle 8



Cycle 9



That's a lot of diagrams there

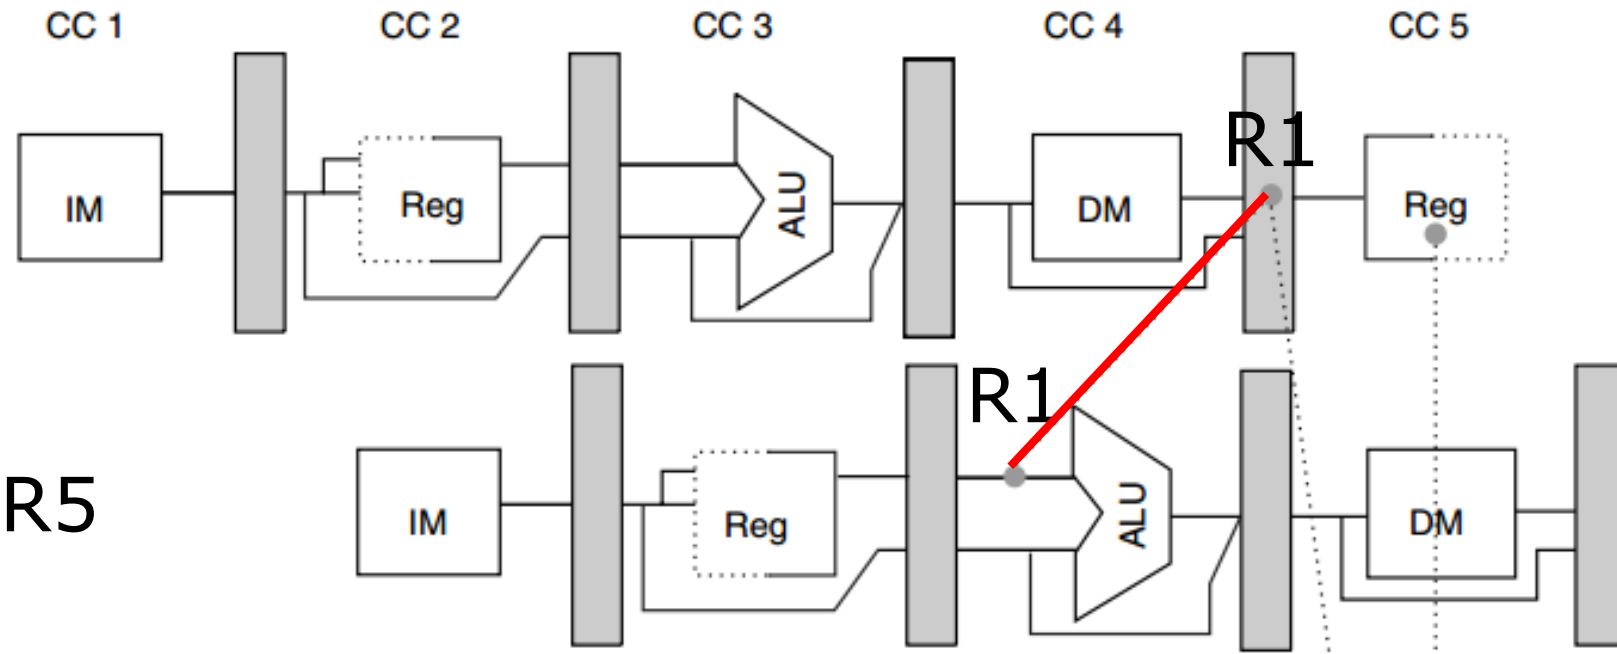
		Clock cycle								
		1	2	3	4	5	6	7	8	9
lw	\$t0, 4(\$sp)	IF	ID	EX	MEM	WB				
sub	\$v0, \$a0, \$a1		IF	ID	EX	MEM	WB			
and	\$t1, \$t2, \$t3			IF	ID	EX	MEM	WB		
or	\$s0, \$s1, \$s2				IF	ID	EX	MEM	WB	
add	\$t5, \$t6, \$0					IF	ID	EX	MEM	WB

- You can see how instruction executions are overlapped
- Each functional unit is used by a *different* instruction in each cycle
- The pipeline registers save control and data values generated in previous clock cycles for later use
- When the pipeline is full in clock cycle 5, all of the hardware units are utilized. This is the ideal situation, and what makes pipelined processors so fast

When Pipeline Is Stuck

LD R1, 0(R2)

SUB R4, R1, R5



Actually, in many cases, instructions cannot be fully pipelined.

In this example, R1 is ready by load instruction at the end of clock cycle 4;

But subtract instruction requires R1 at the beginning of clock cycle 4; in this case, the processor can not fully pipeline the two instructions as we expect.

References

- Computer System Architecture (3rd Edition) by M. Morris Mano
- Computer organization and architecture designing for performance (8th edition) by William Stallings
- Computer Architecture: A Quantitative Approach 5th Edition by John L. Hennessy
- Patterson & Hennessy, “Computer Organization & Design”
- Lecture Notes, Lecture 4: Pipelining Basics & Hazards, Kai Bu, kaibu@zju.edu.cn
- EmbeededNotes Pipeline ,
https://www.mediafire.com/file/yl157ng7eby7989/Video_6_-_PipeLine_-_1_Intro.pptx/file
- <https://www.youtube.com/watch?v=X84uO3H83bA>
- Lecture Notes, Introduction to Computer Architecture, at the University of California, Santa Barbara. © Behrooz Parhami

Thanks

